

# Making Predictions

PREDICTING CUSTOMER CHURN IN PYTHON



**Mark Peterson**

Senior Data Scientist, Alliance Data

# (Supervised) Machine Learning Primer

- Goal: Predict whether or not a customer will churn
- Target Variable: 'Churn'
- Supervised Machine Learning
- Learn from historical (training) data to make new predictions



# Model Selection

- Which model to use?
- ... it depends!
- In this course: Experiment with several models
- To learn about their inner workings: Check out other DataCamp courses

# Model Selection

- Logistic regression: Good baseline
  - Offers simplicity and interpretability
  - Cannot capture more complex relationships
- Random forests
- Support vector machines

# Training your Model

```
from sklearn.svm import SVC

svc = SVC()

svc.fit(telco[features], telco['target'])
```

```
SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
    decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
    max_iter=-1, probability=False, random_state=None, shrinking=True,
    tol=0.001, verbose=False)
```

# Making a Prediction

```
prediction = svc.predict(new_customer)

print(prediction)
```

```
[0]
```

# Let's practice!

PREDICTING CUSTOMER CHURN IN PYTHON



# Evaluating Model Performance

PREDICTING CUSTOMER CHURN IN PYTHON



**Mark Peterson**

Senior Data Scientist, Alliance Data

# Accuracy

- One possible metric: Accuracy
  - $\text{Total Number of Correct Predictions} / \text{Total Number of Data Points}$
- What data to use?
  - Training data not representative of new data

# Training and Test Sets

- Fit your classifier to the training set
- Make predictions using the test set

# Training and Test Sets using scikit-learn

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(telco['data'], telco['target'],
                                                    test_size=0.2, random_state = 42)

from sklearn.svm import SVC

svc = SVC()

svc.fit(X_train, y_train)

svc.predict(X_test)
```

# Computing Accuracy

```
svc.score(X_test, y_test)
```

```
0.857
```

- 85.7% accuracy: Quite good for a first try!

# Improving your model

- Overfitting: Model fits the training data too closely
- Underfitting: Does not capture trends in the training data
- Need to find the right balance between overfitting and underfitting

# Let's practice!

PREDICTING CUSTOMER CHURN IN PYTHON

# Model Metrics

PREDICTING CUSTOMER CHURN IN PYTHON



**Mark Peterson**

Senior Data Scientist, Alliance Data



# Imbalanced classes

```
telco['Churn'].value_counts()
```

```
no    2850  
yes    483  
Name: Churn, dtype: int64
```

- Accuracy not a very useful metric

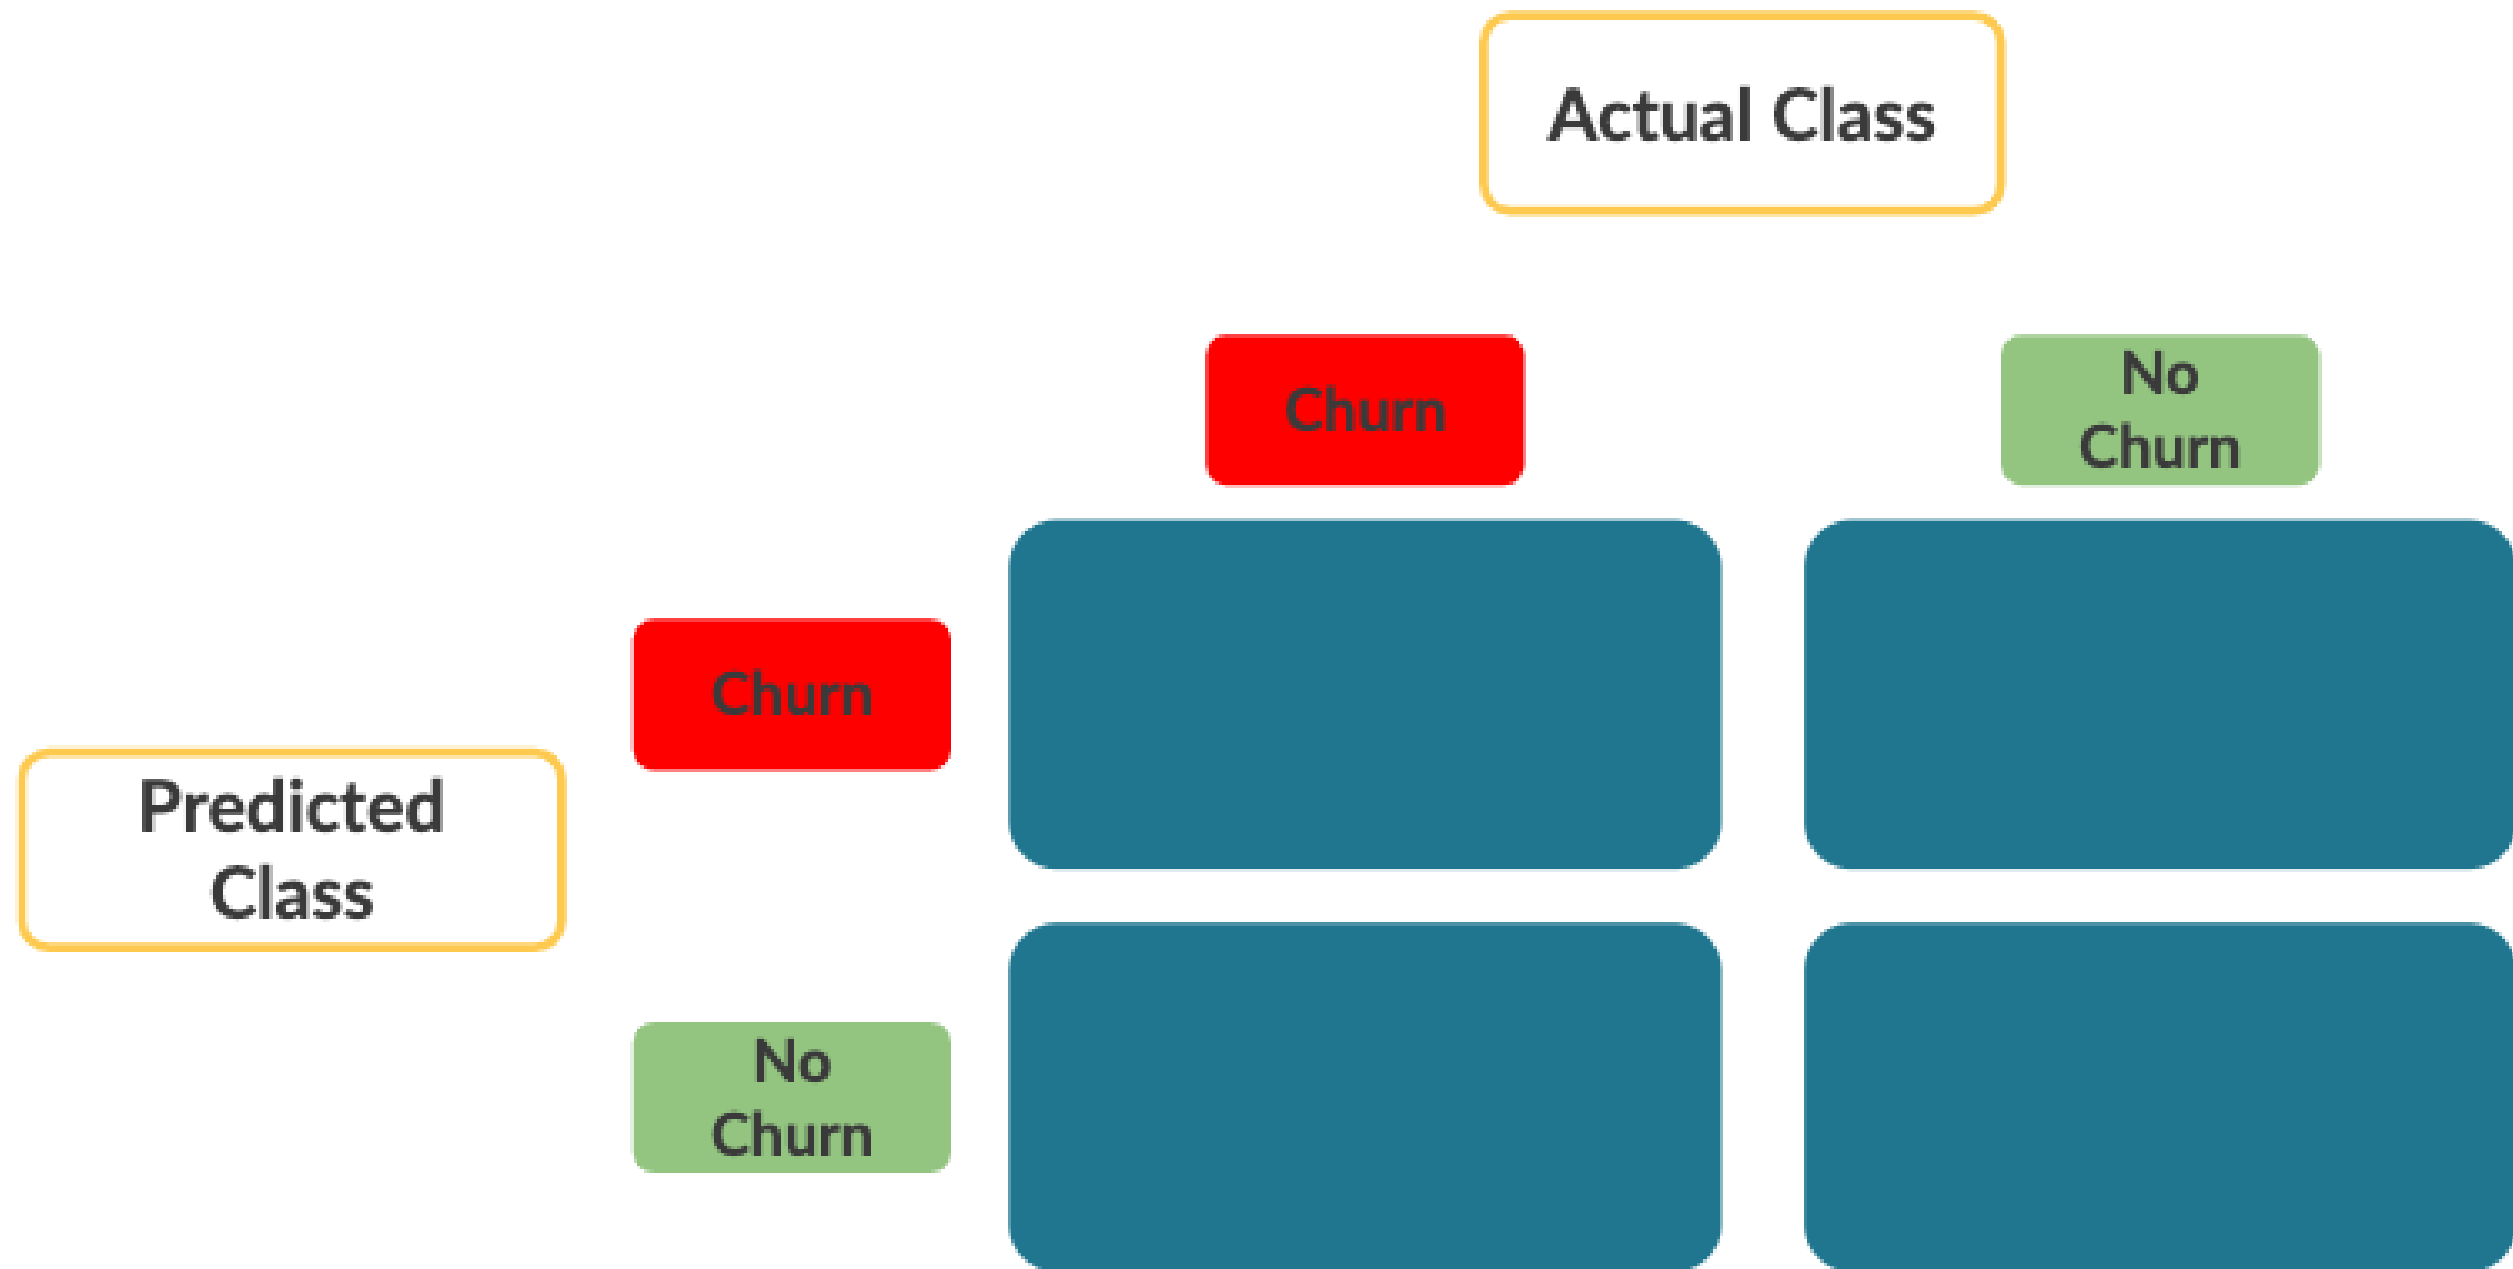


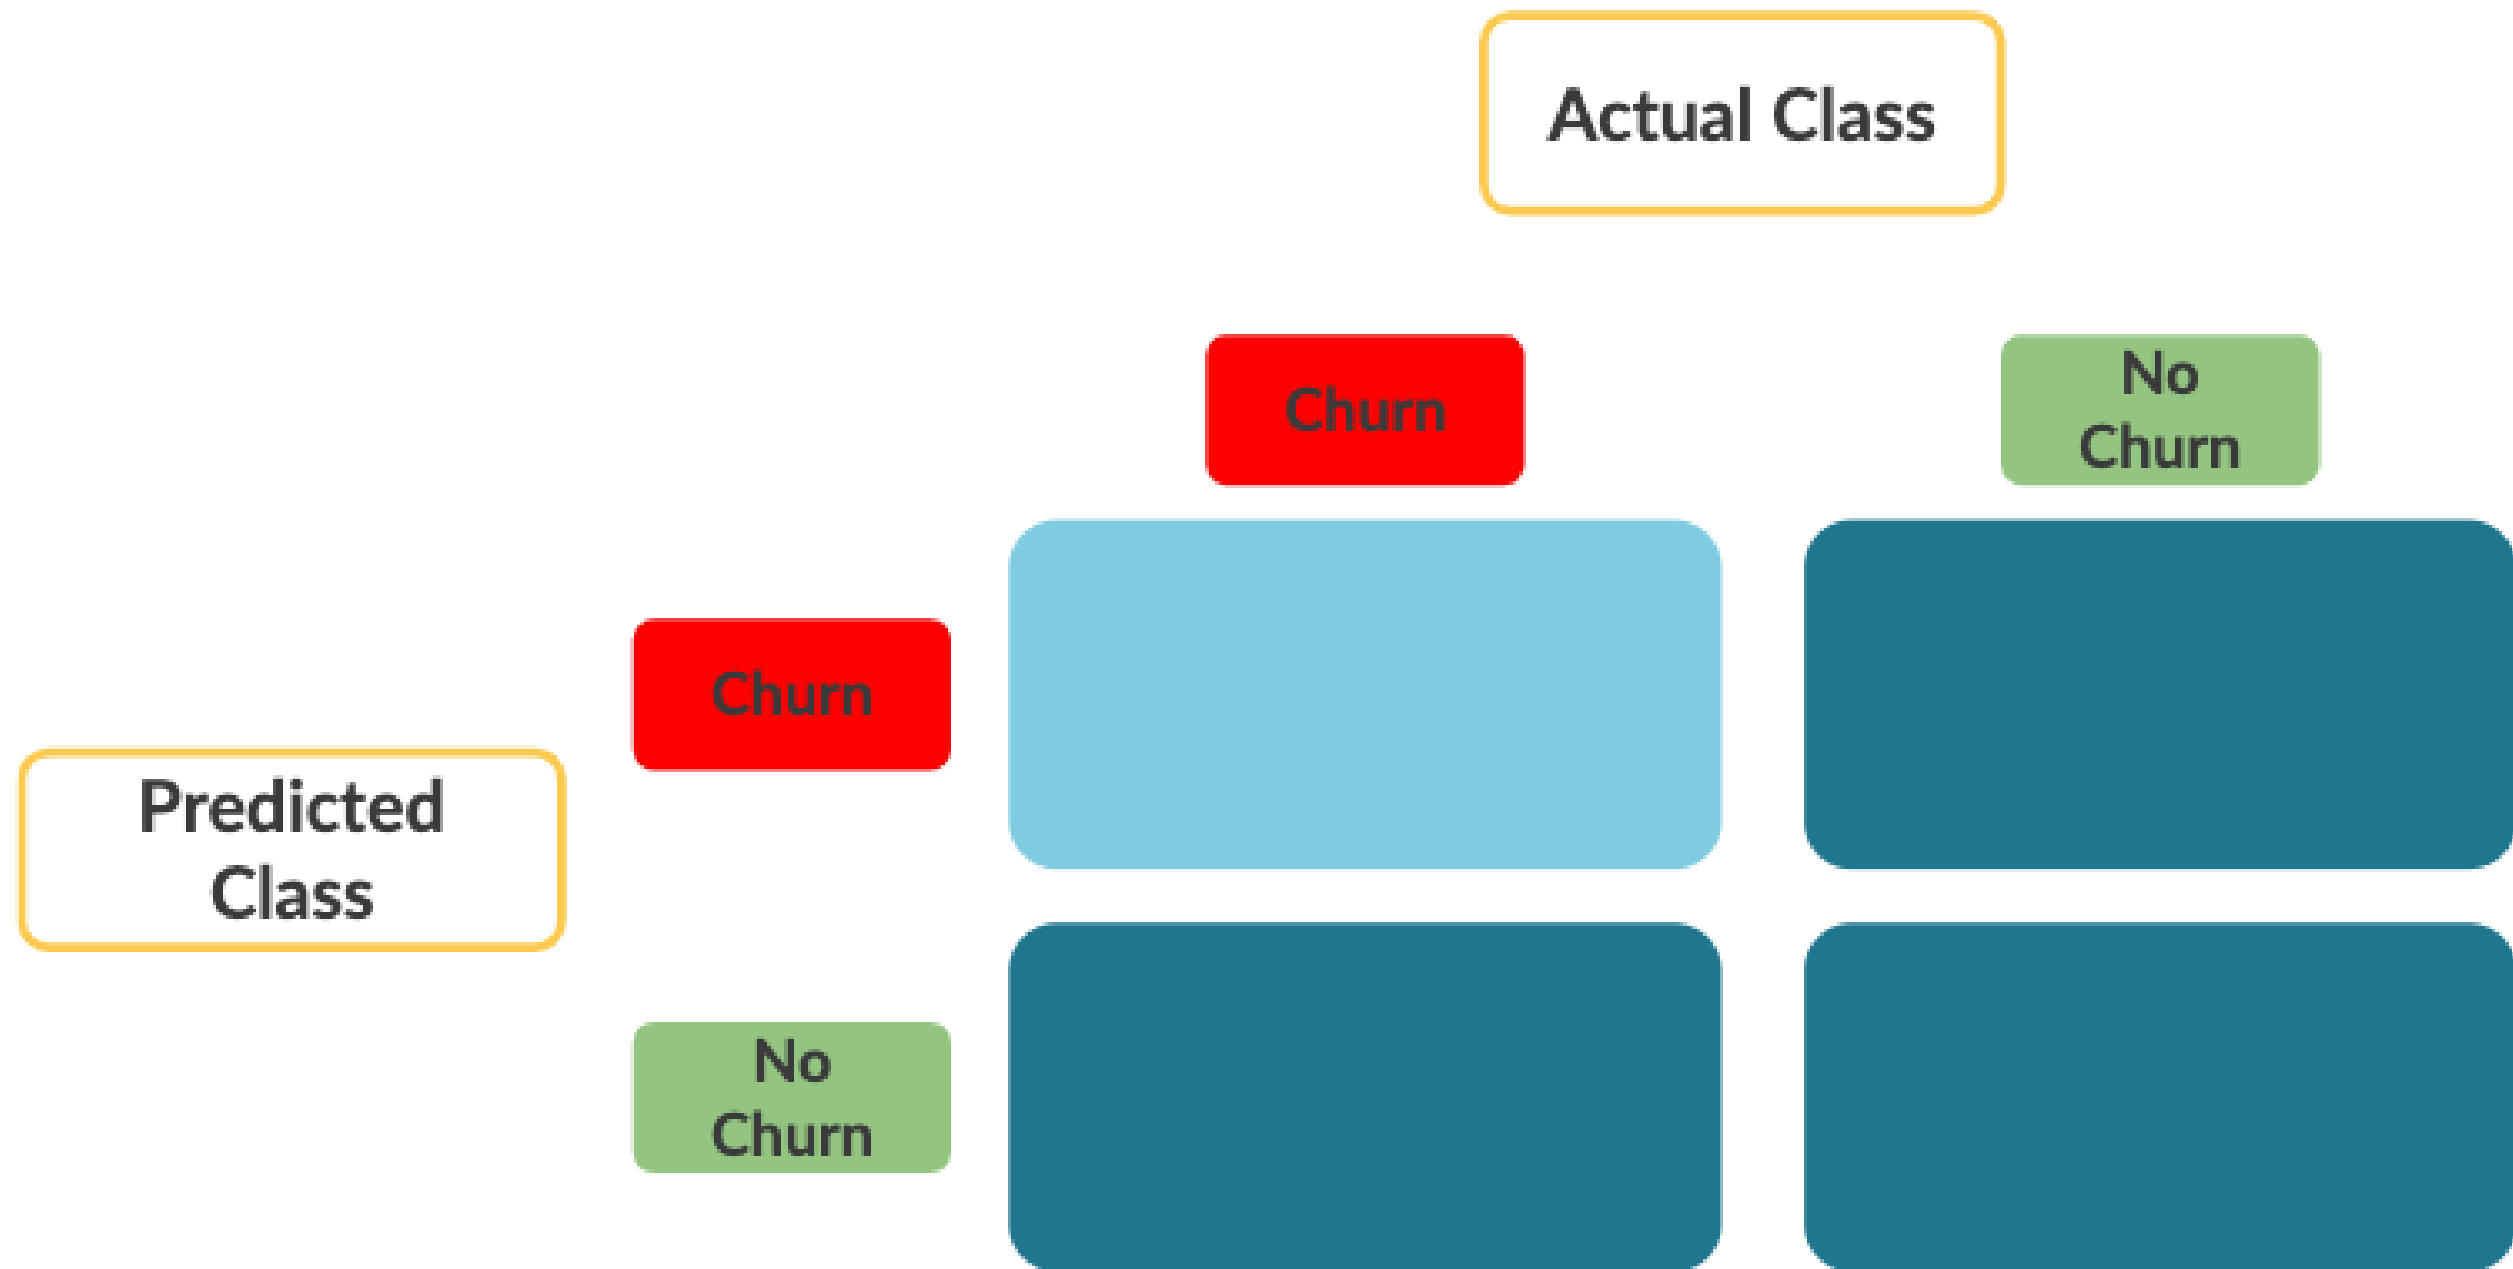
Actual Class

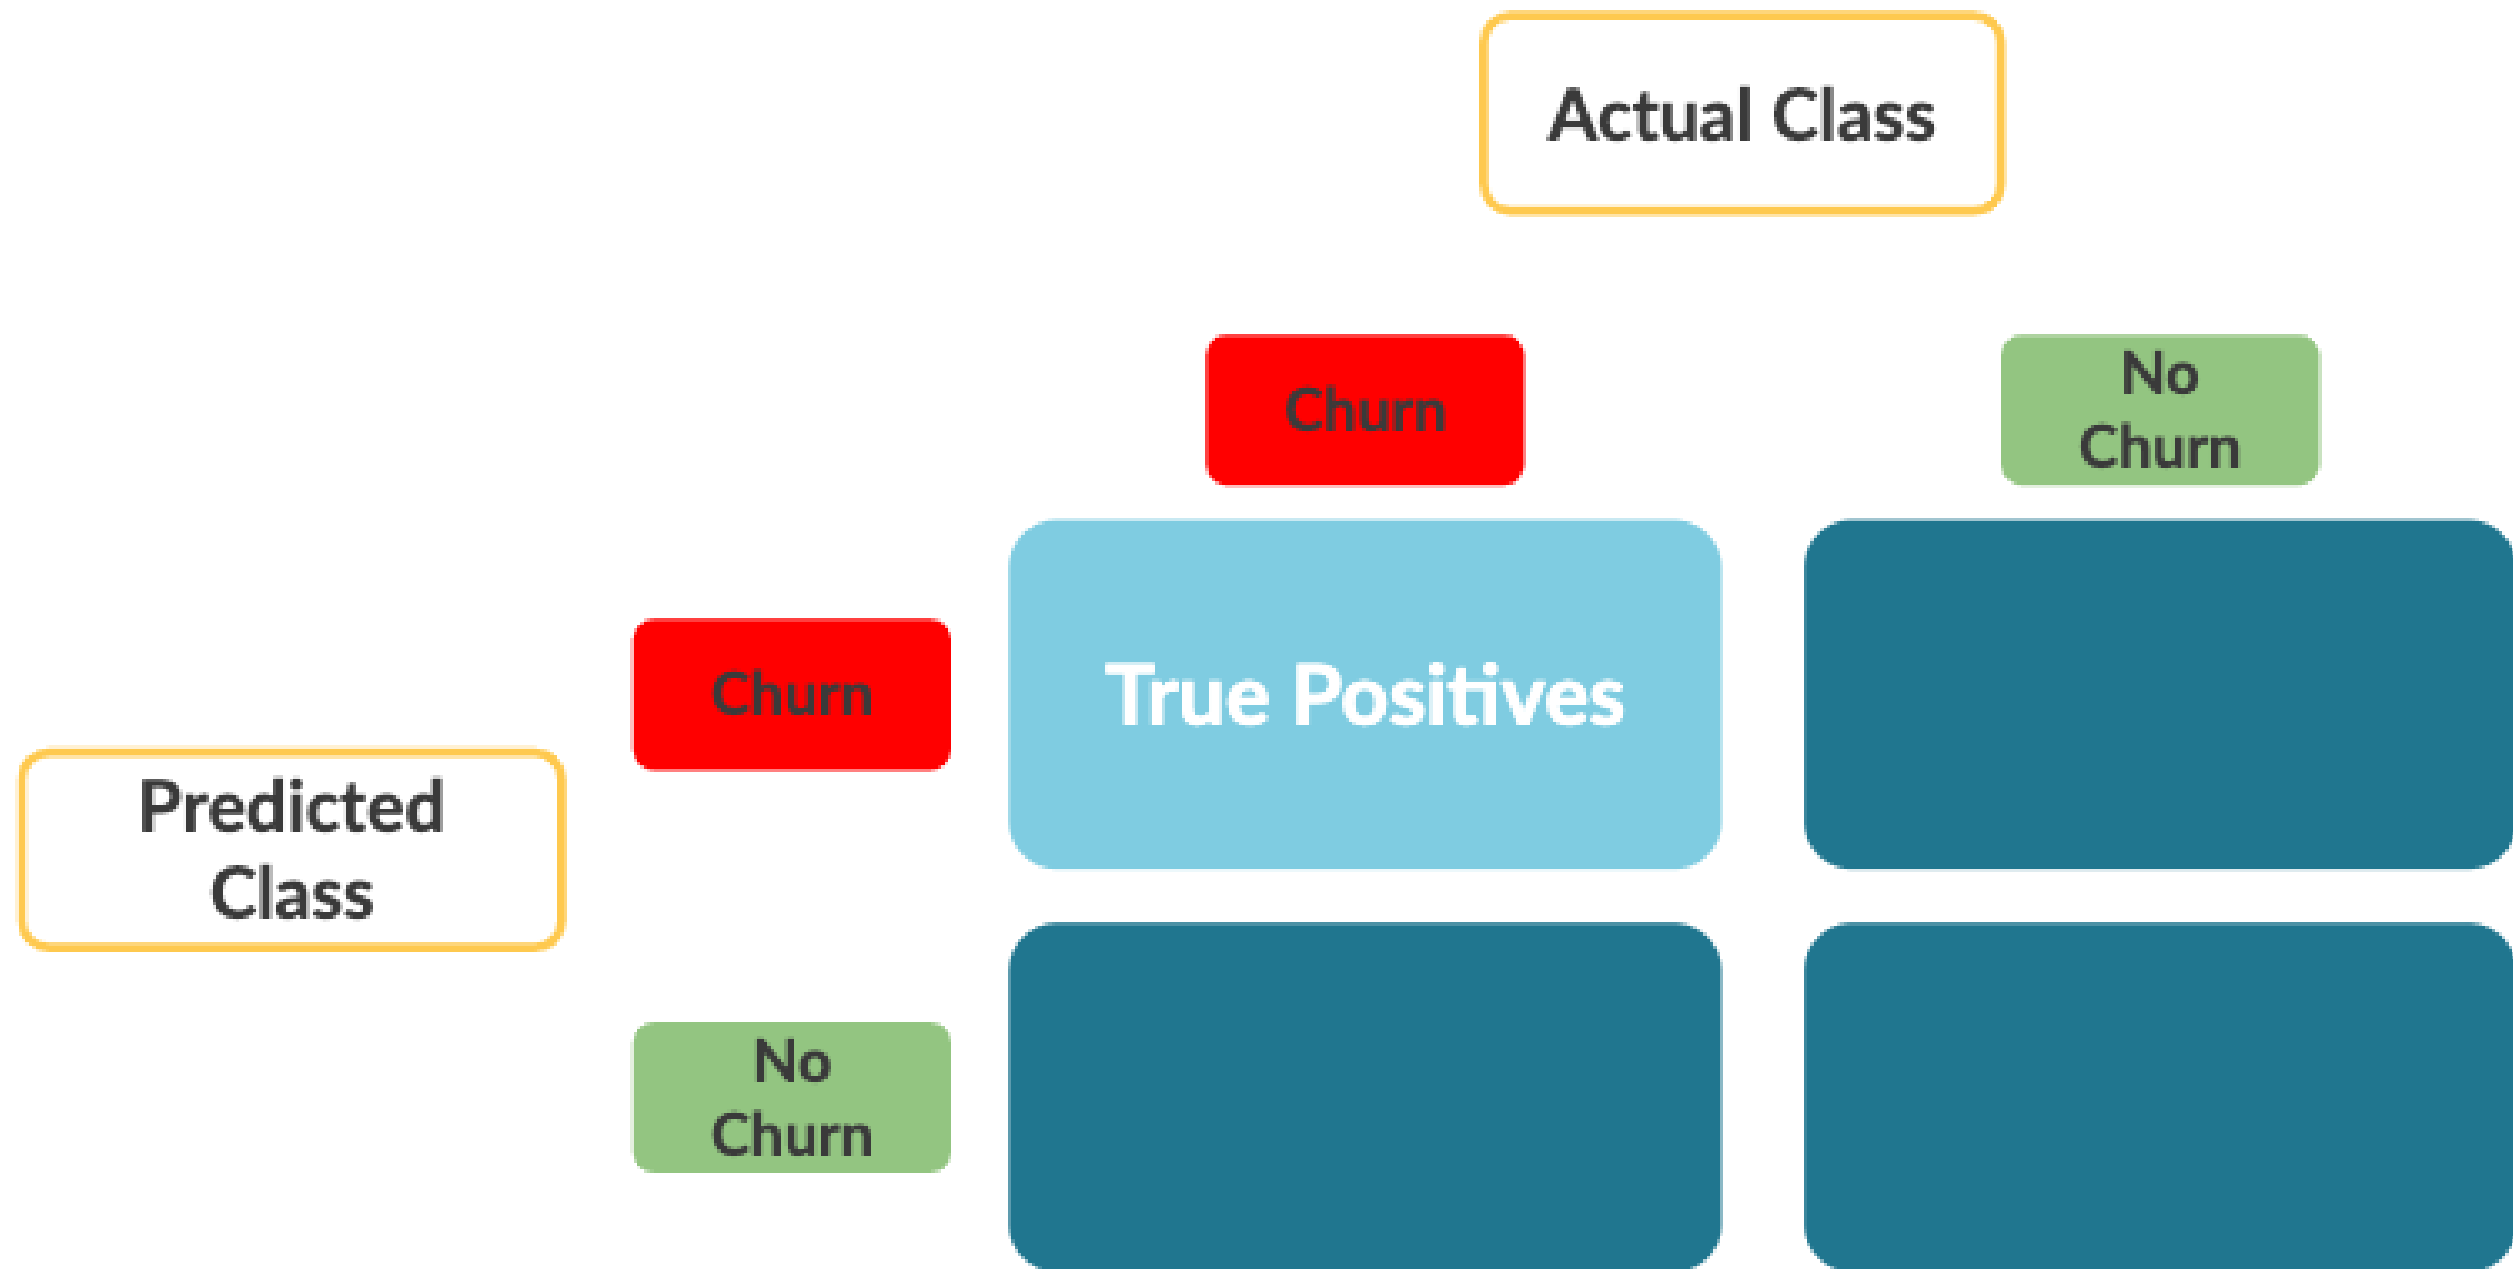
Churn

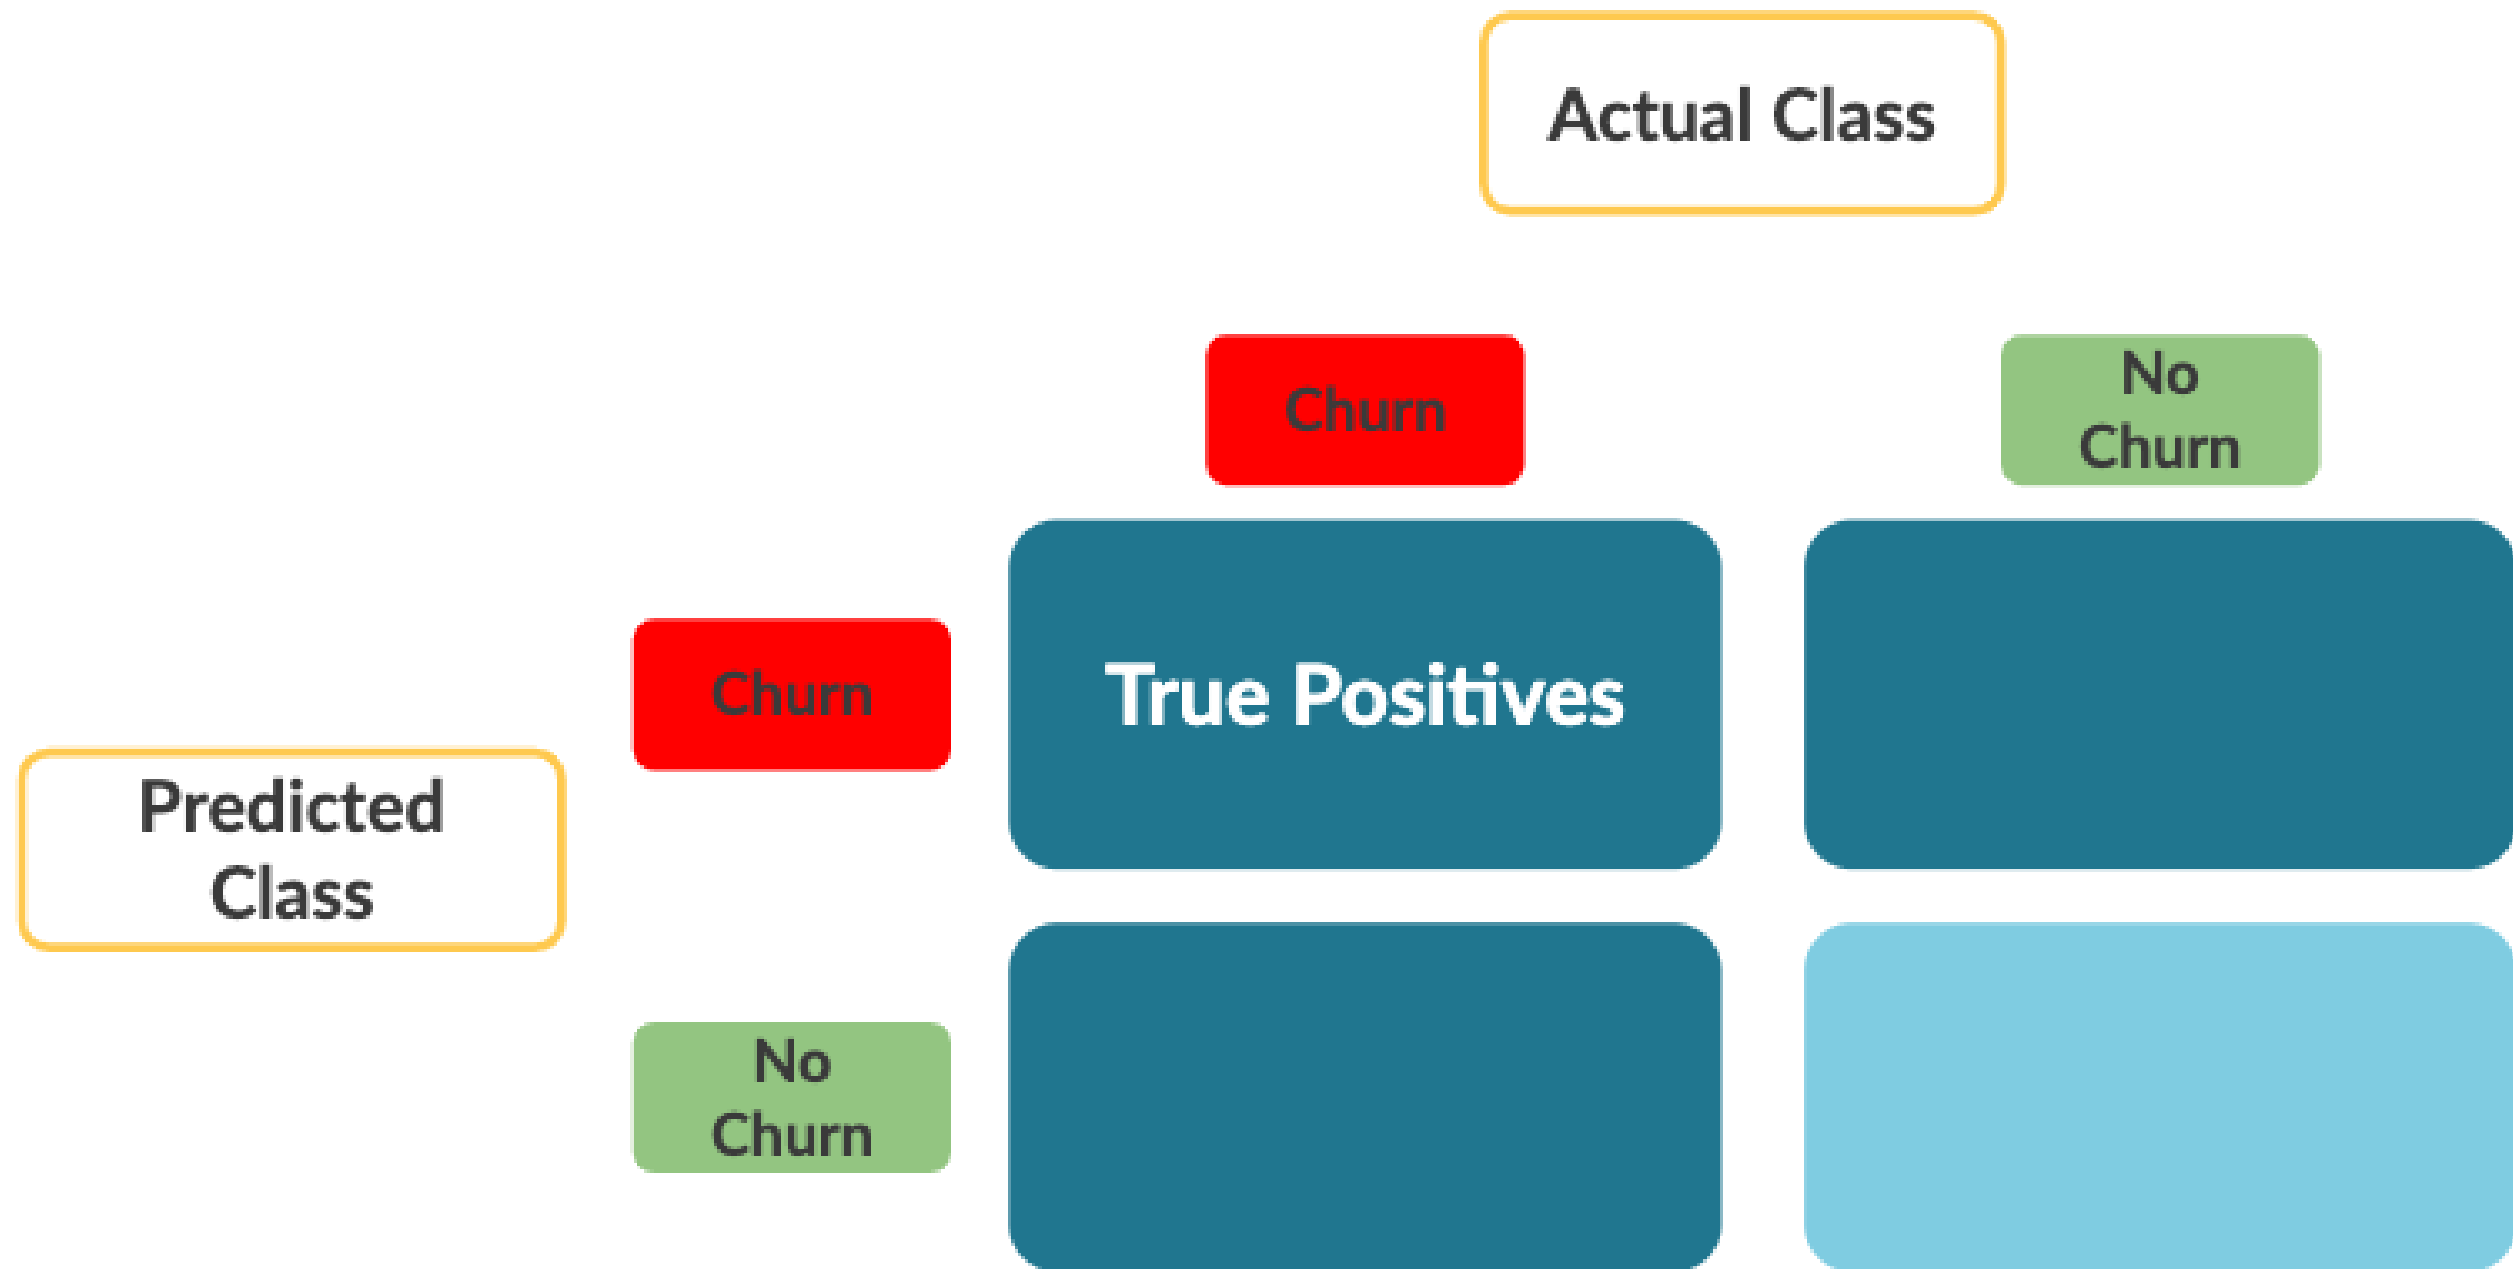
No  
Churn

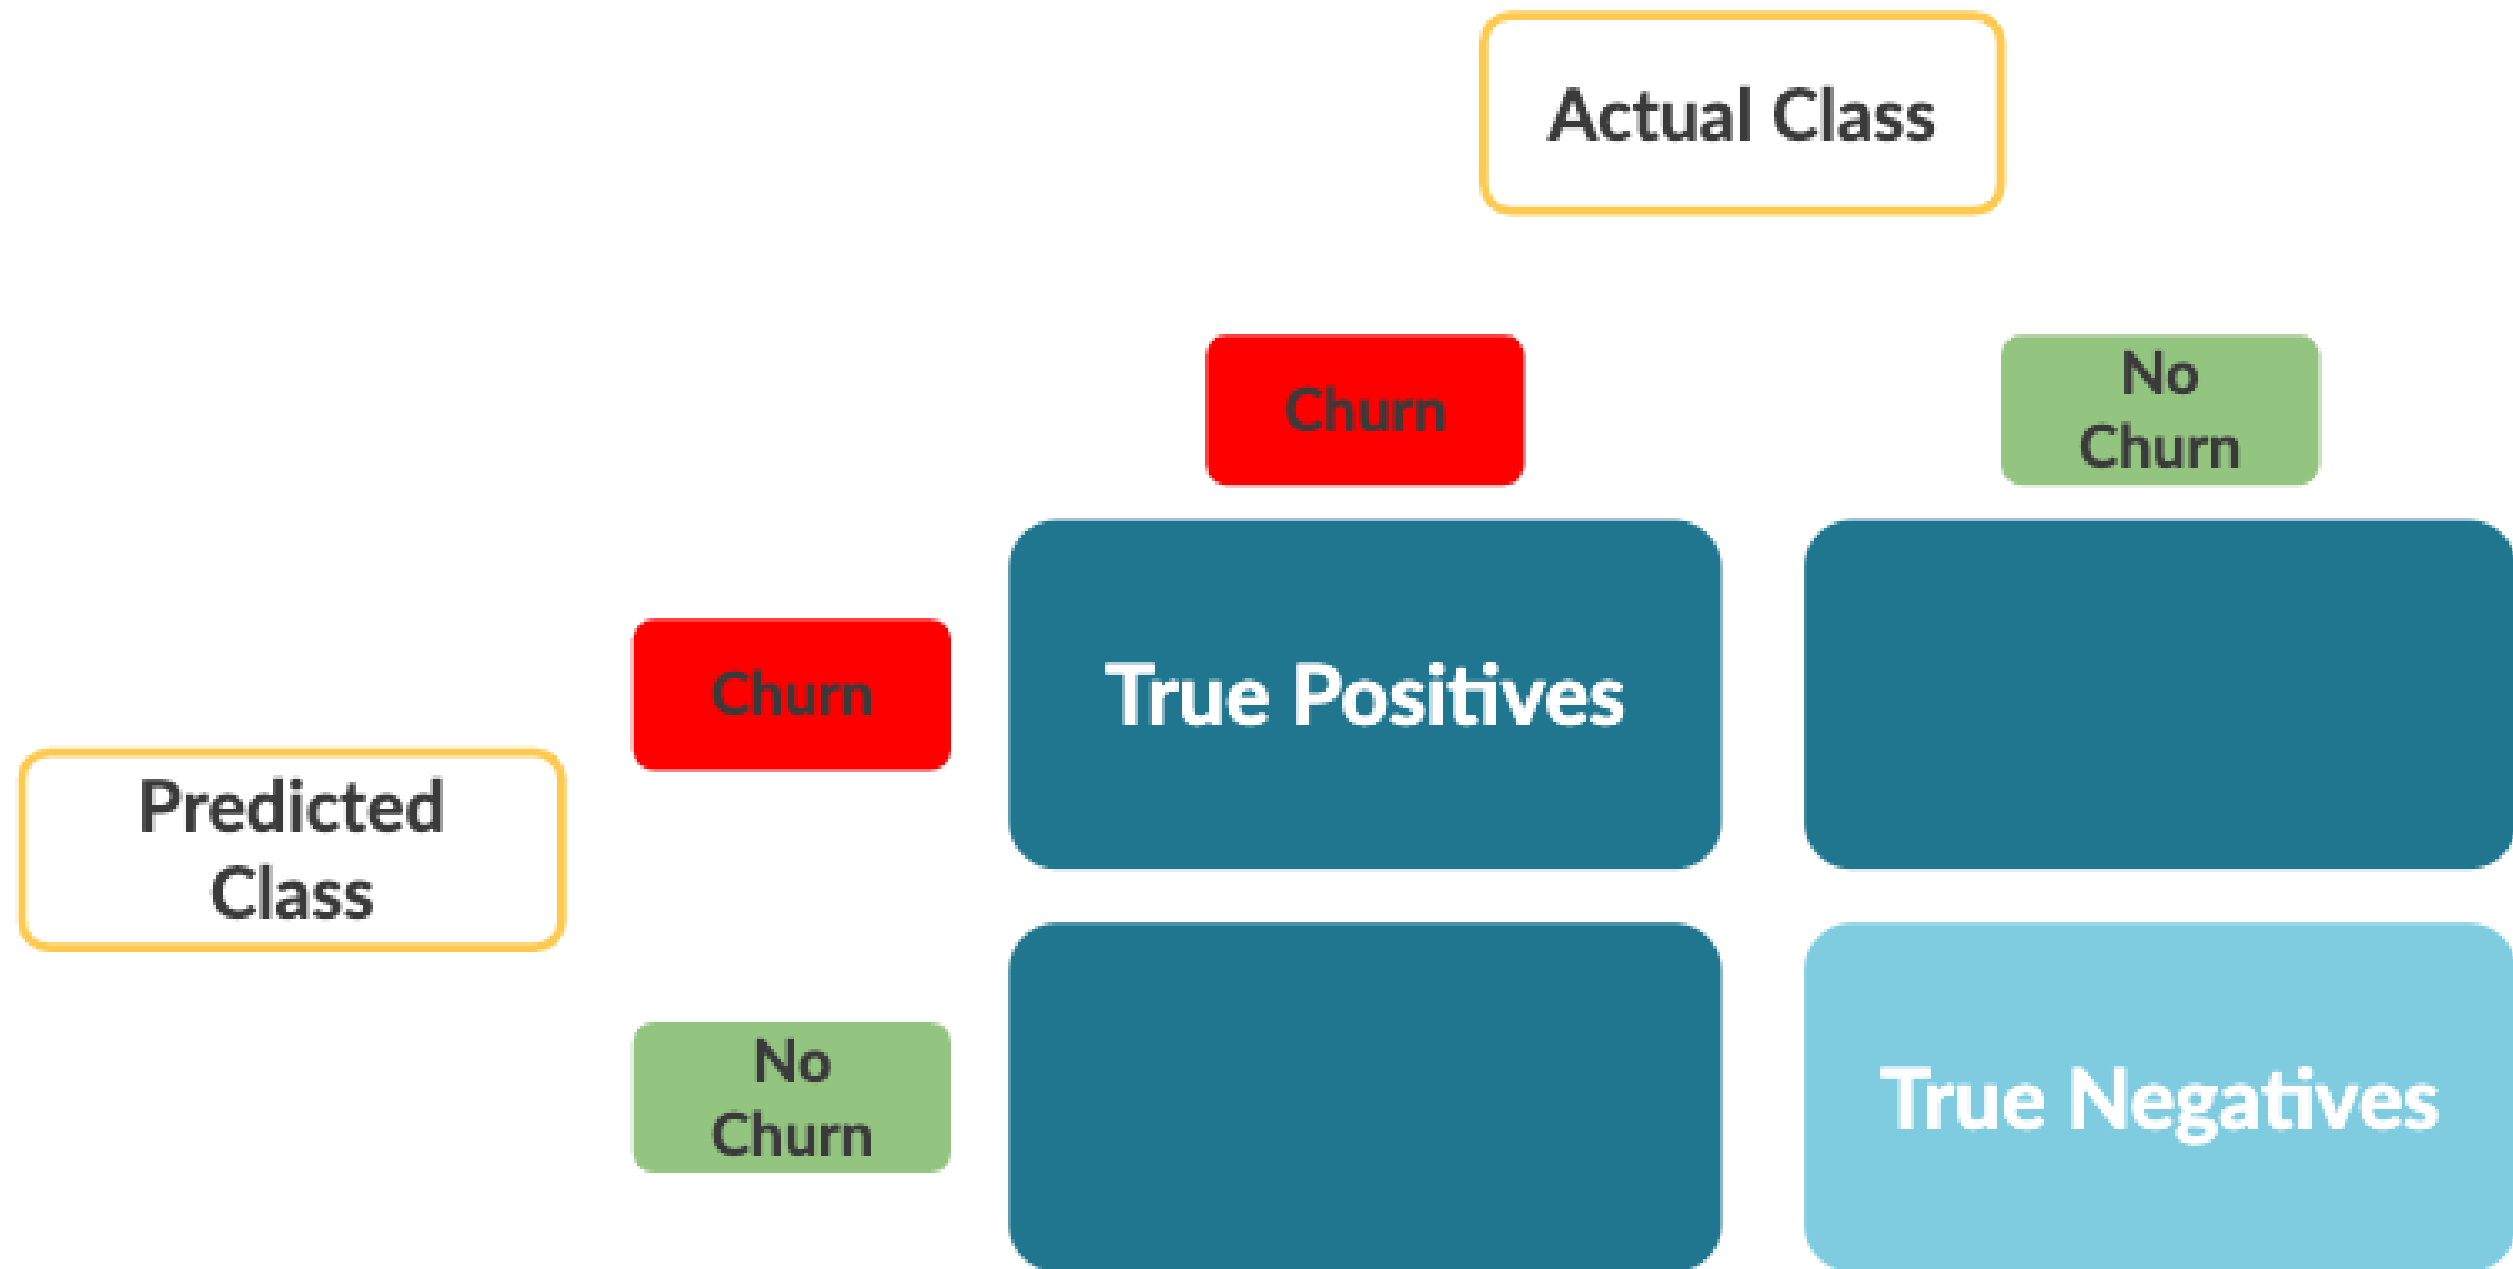




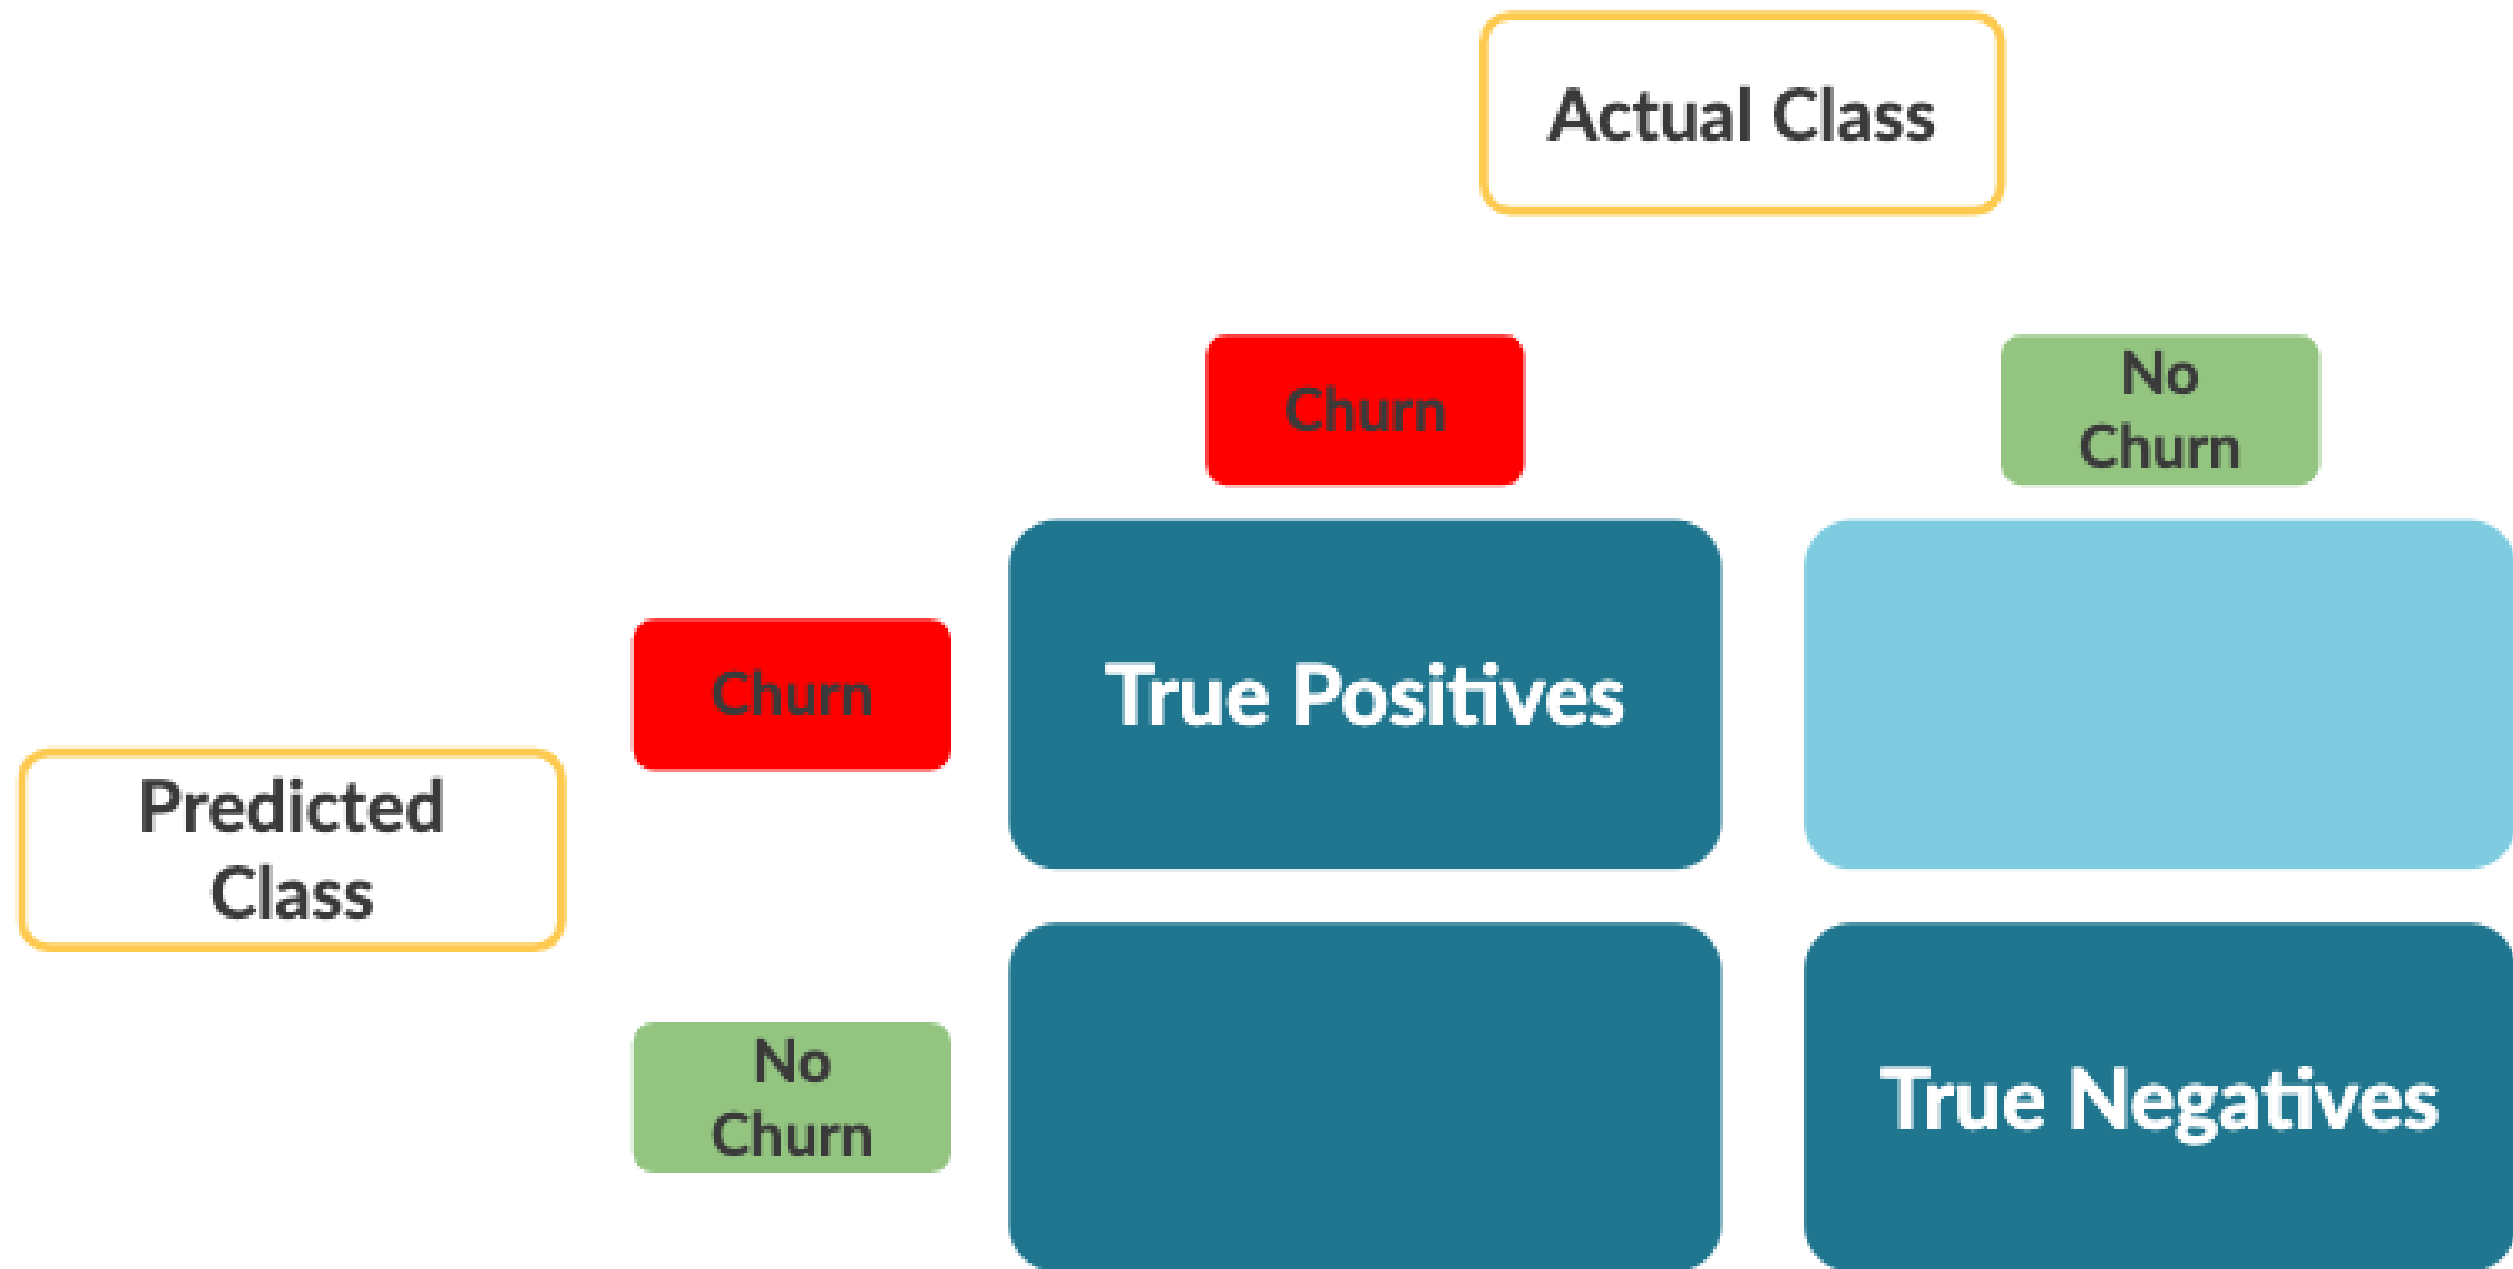


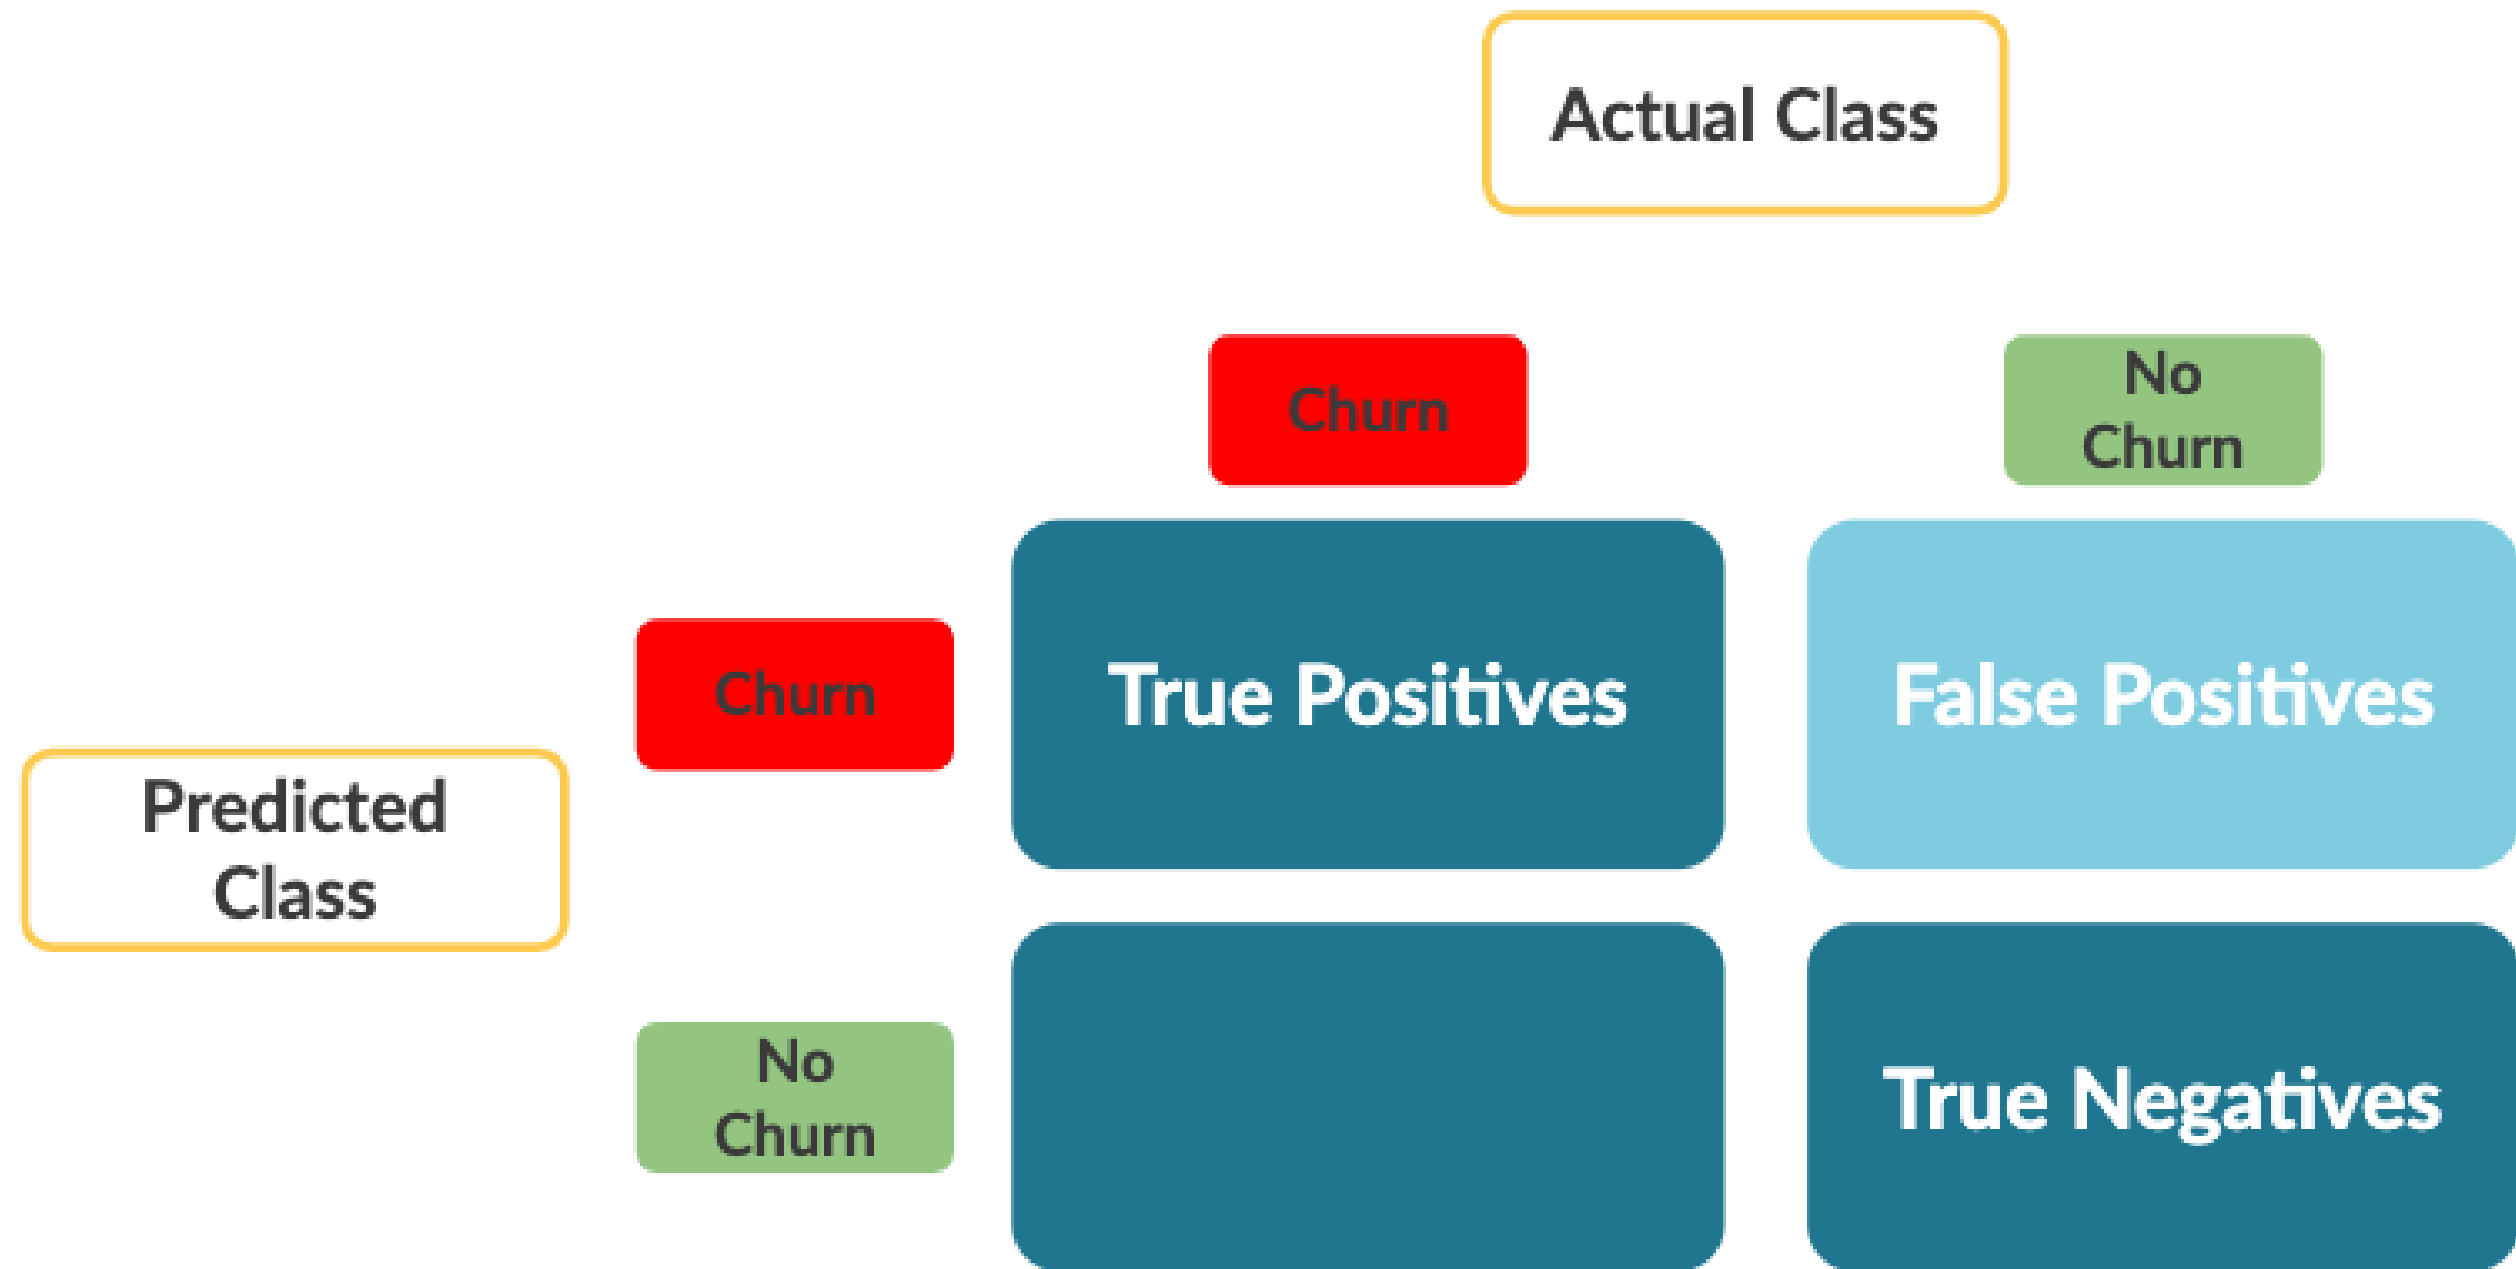


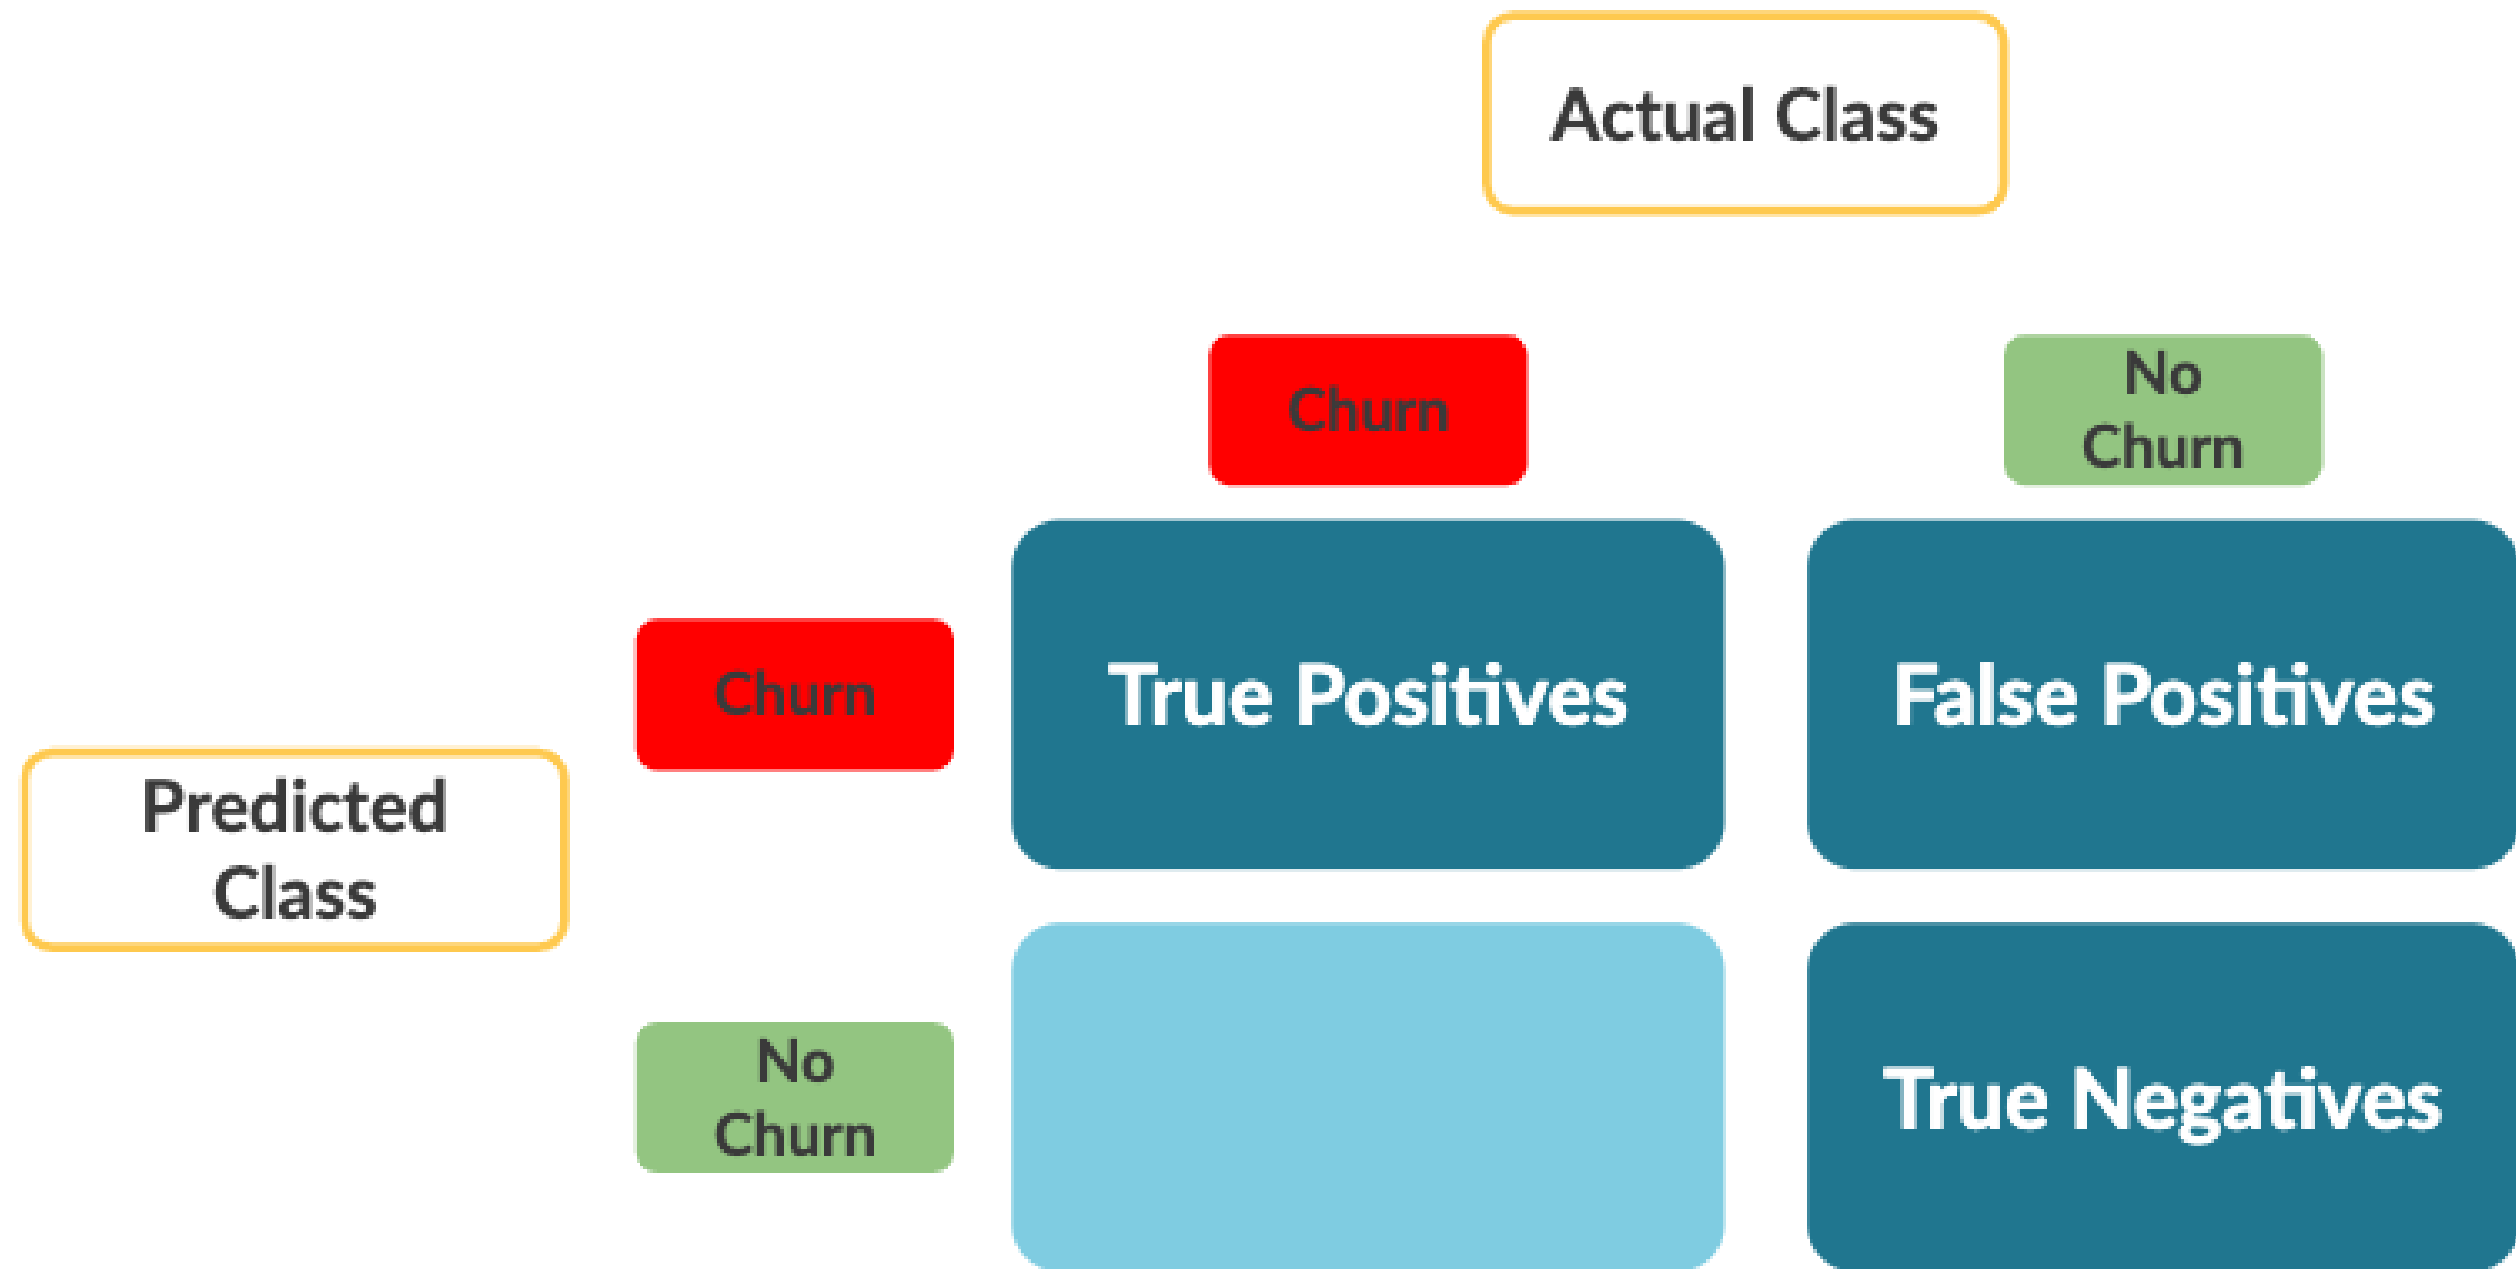


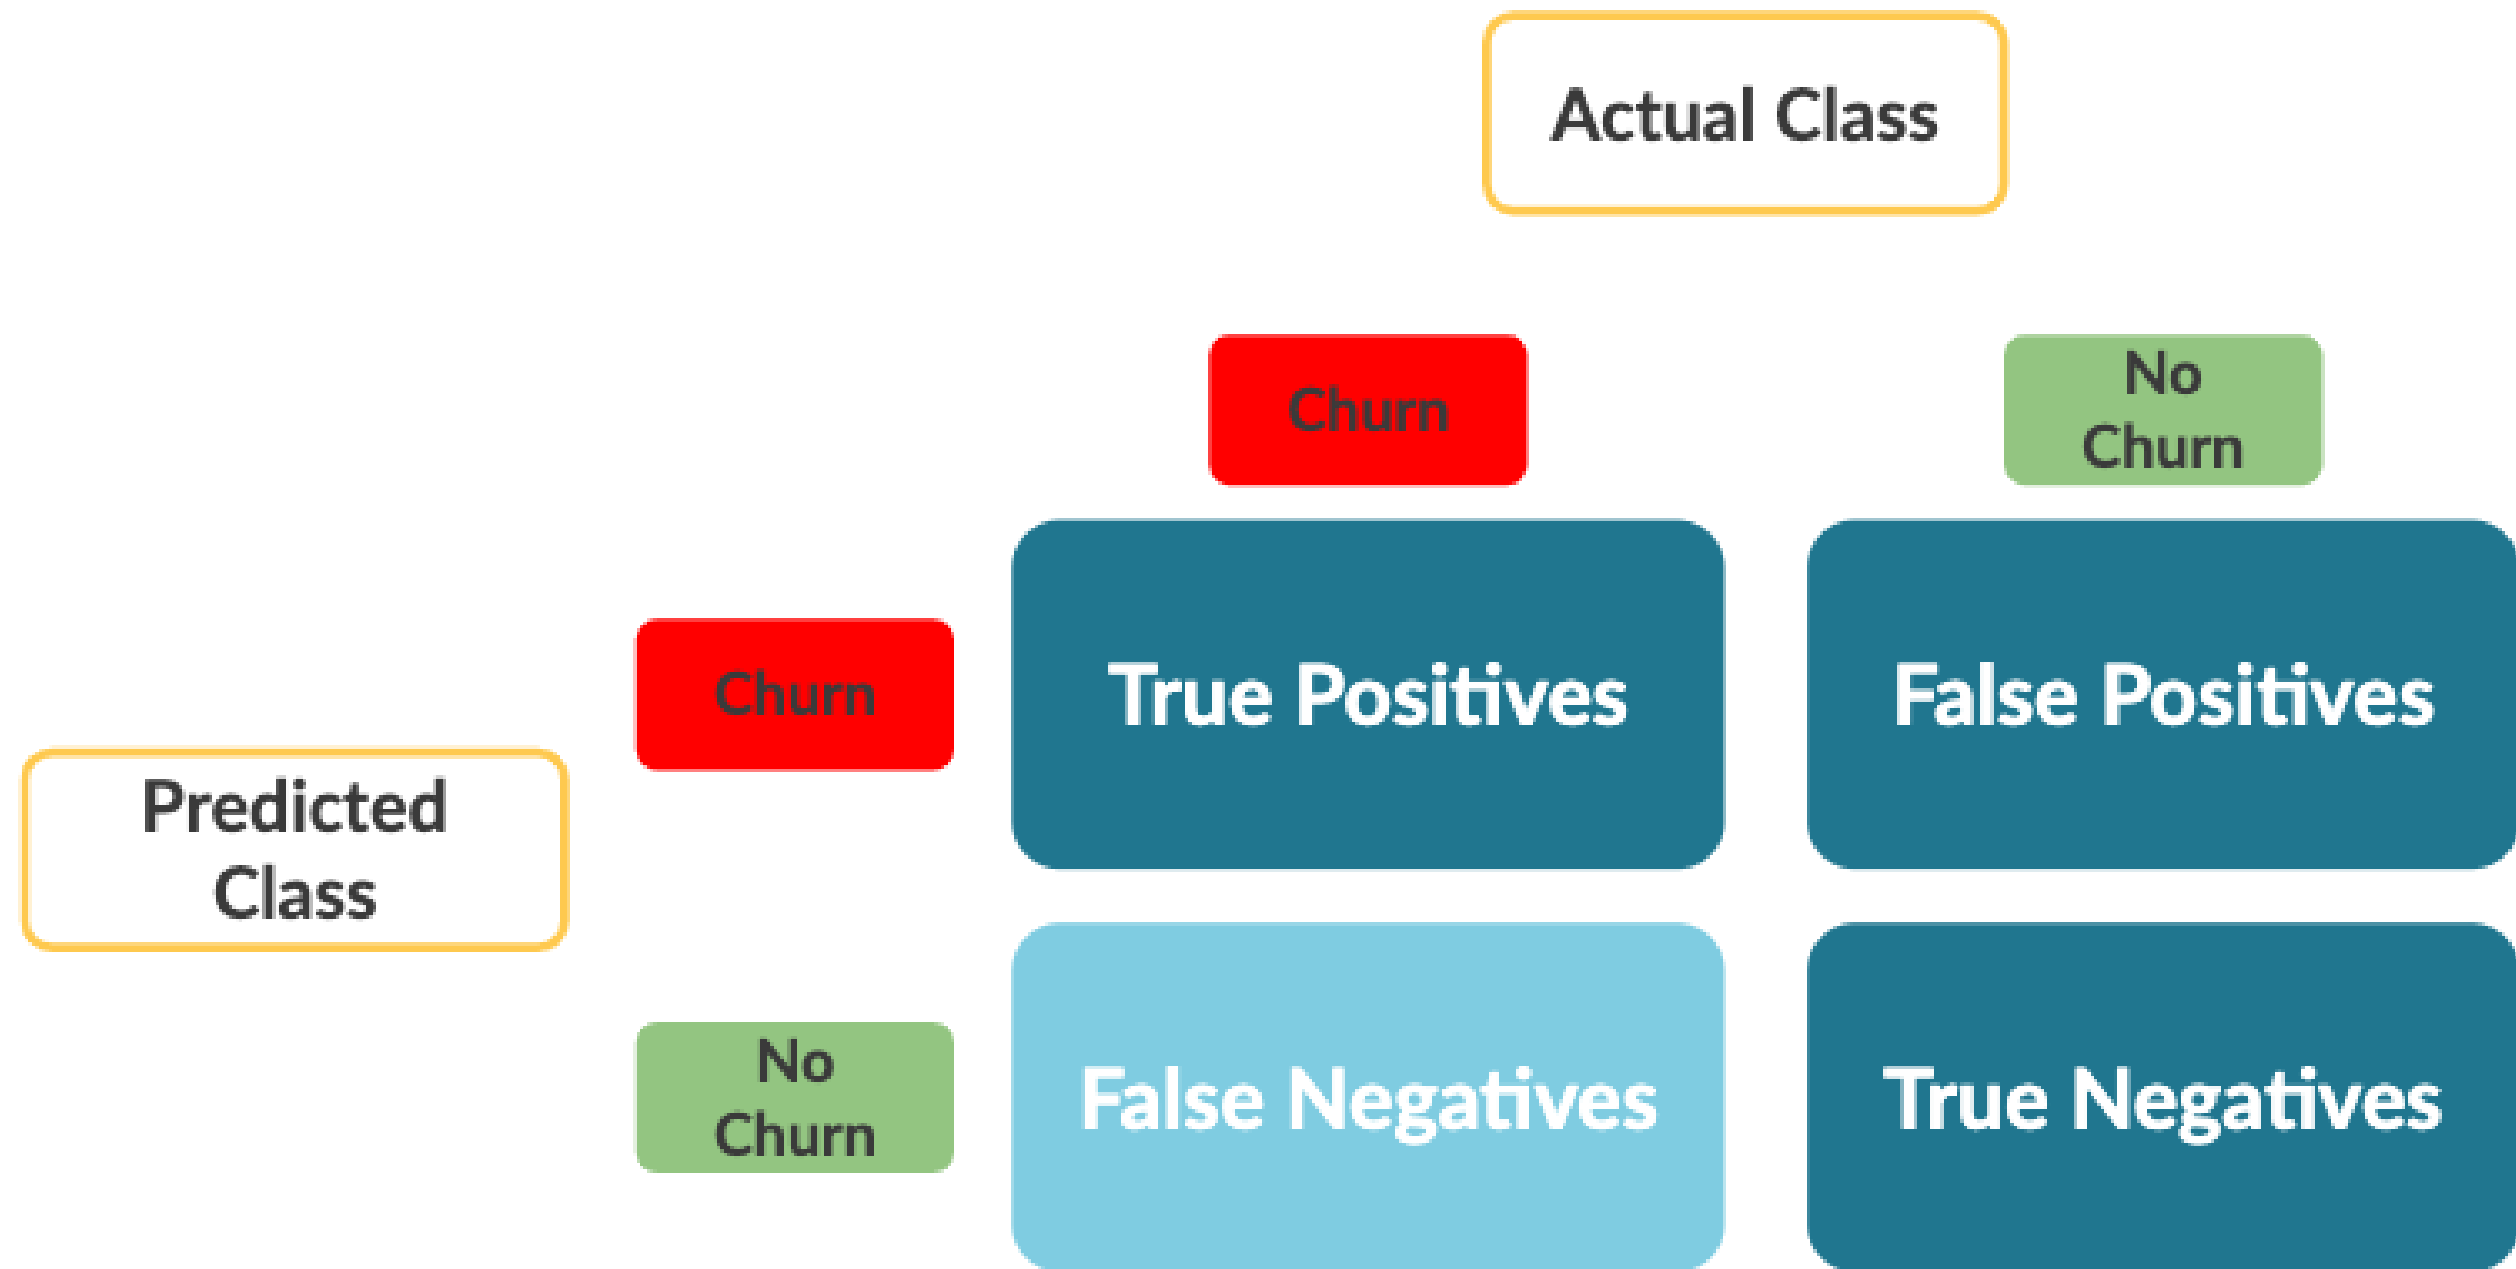


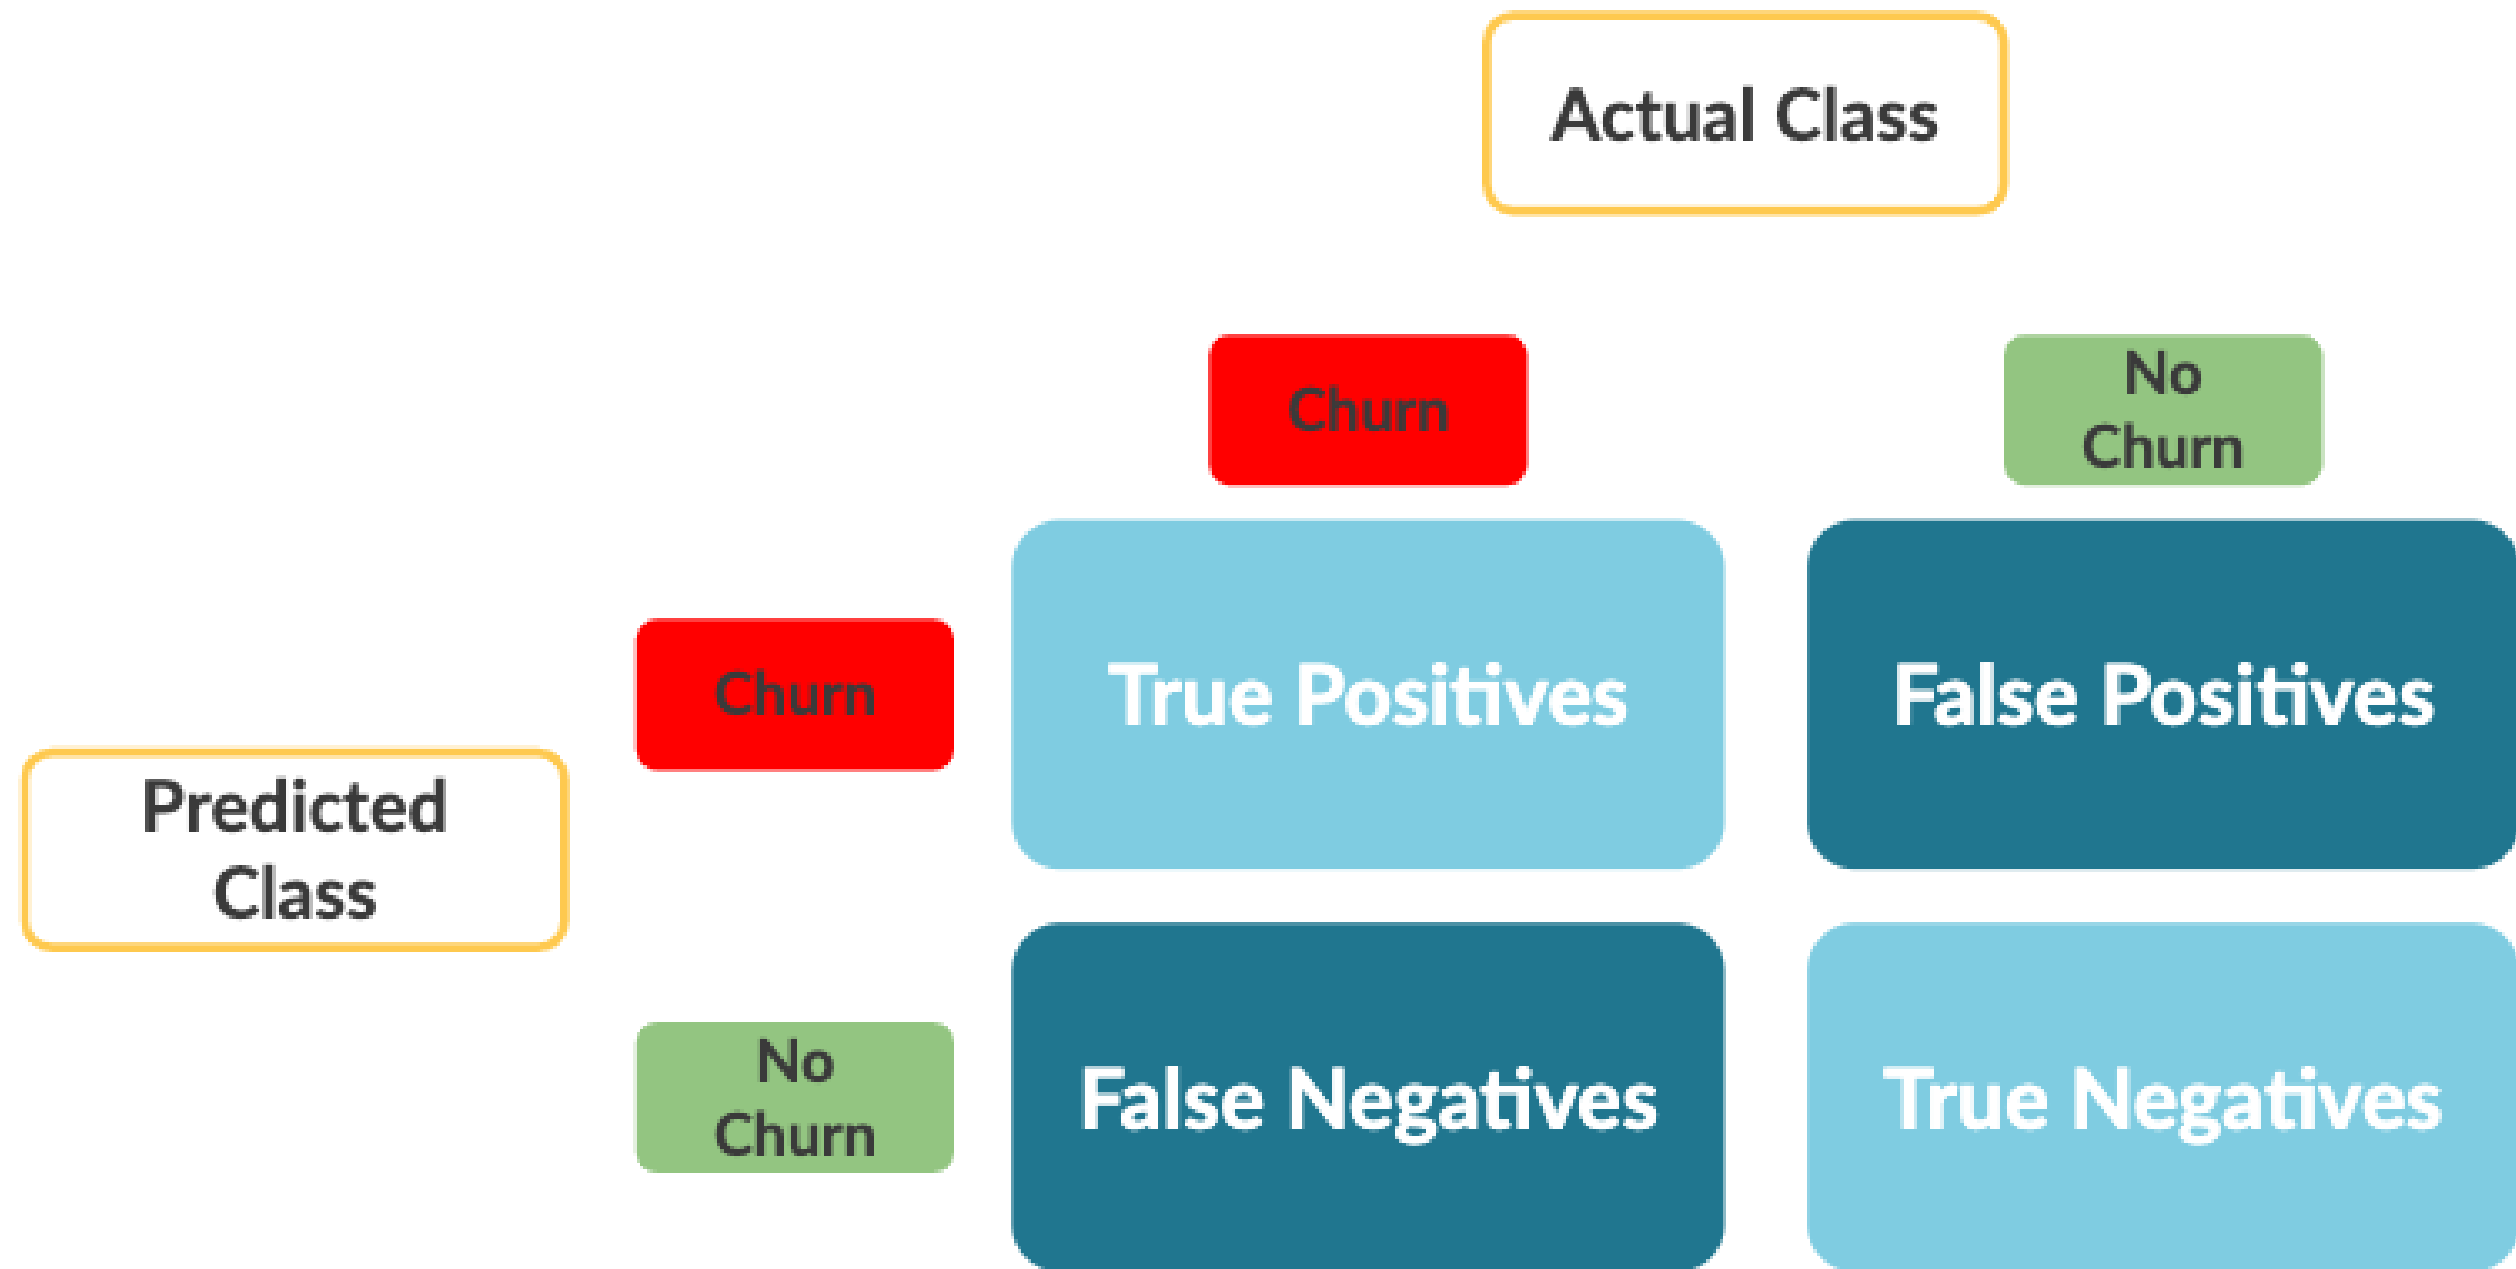












# Precision

Metric	Formula	
Precision	$\text{True Positives} / (\text{True Positives} + \text{False Positives})$	

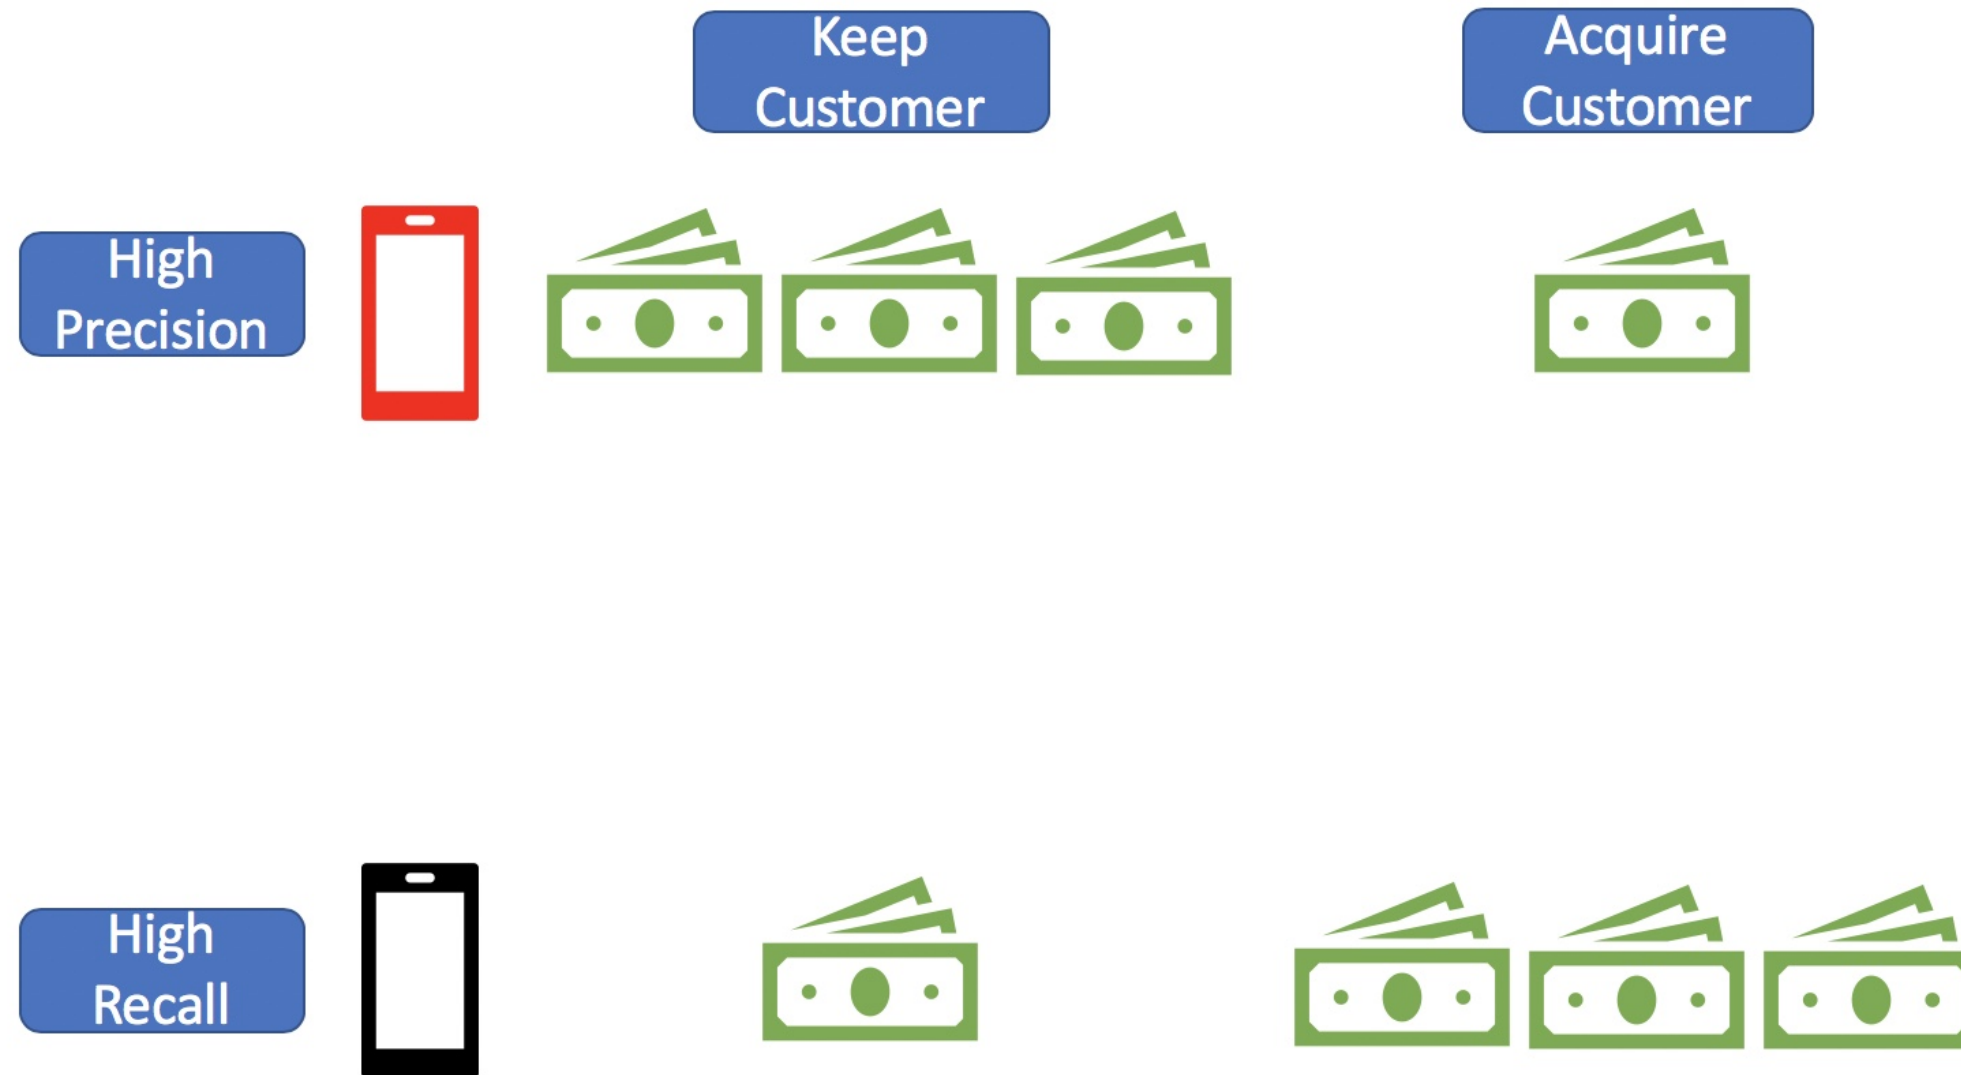
- A model with high precision indicates:
  - Few false positives ("false alarms")
  - Not many non-churners were classified as churners

# Recall

Metric	Formula	
Recall/Sensitivity	$\text{True Positives} / (\text{True Positives} + \text{False Negatives})$	

- A model with high recall indicates that it correctly classified most churners

# Precision vs. Recall





# Confusion Matrix in scikit-learn

```
from sklearn.metrics import confusion_matrix  
  
cm = confusion_matrix(y_test, y_pred)
```

# Let's practice!

PREDICTING CUSTOMER CHURN IN PYTHON

# Other model metrics

PREDICTING CUSTOMER CHURN IN PYTHON



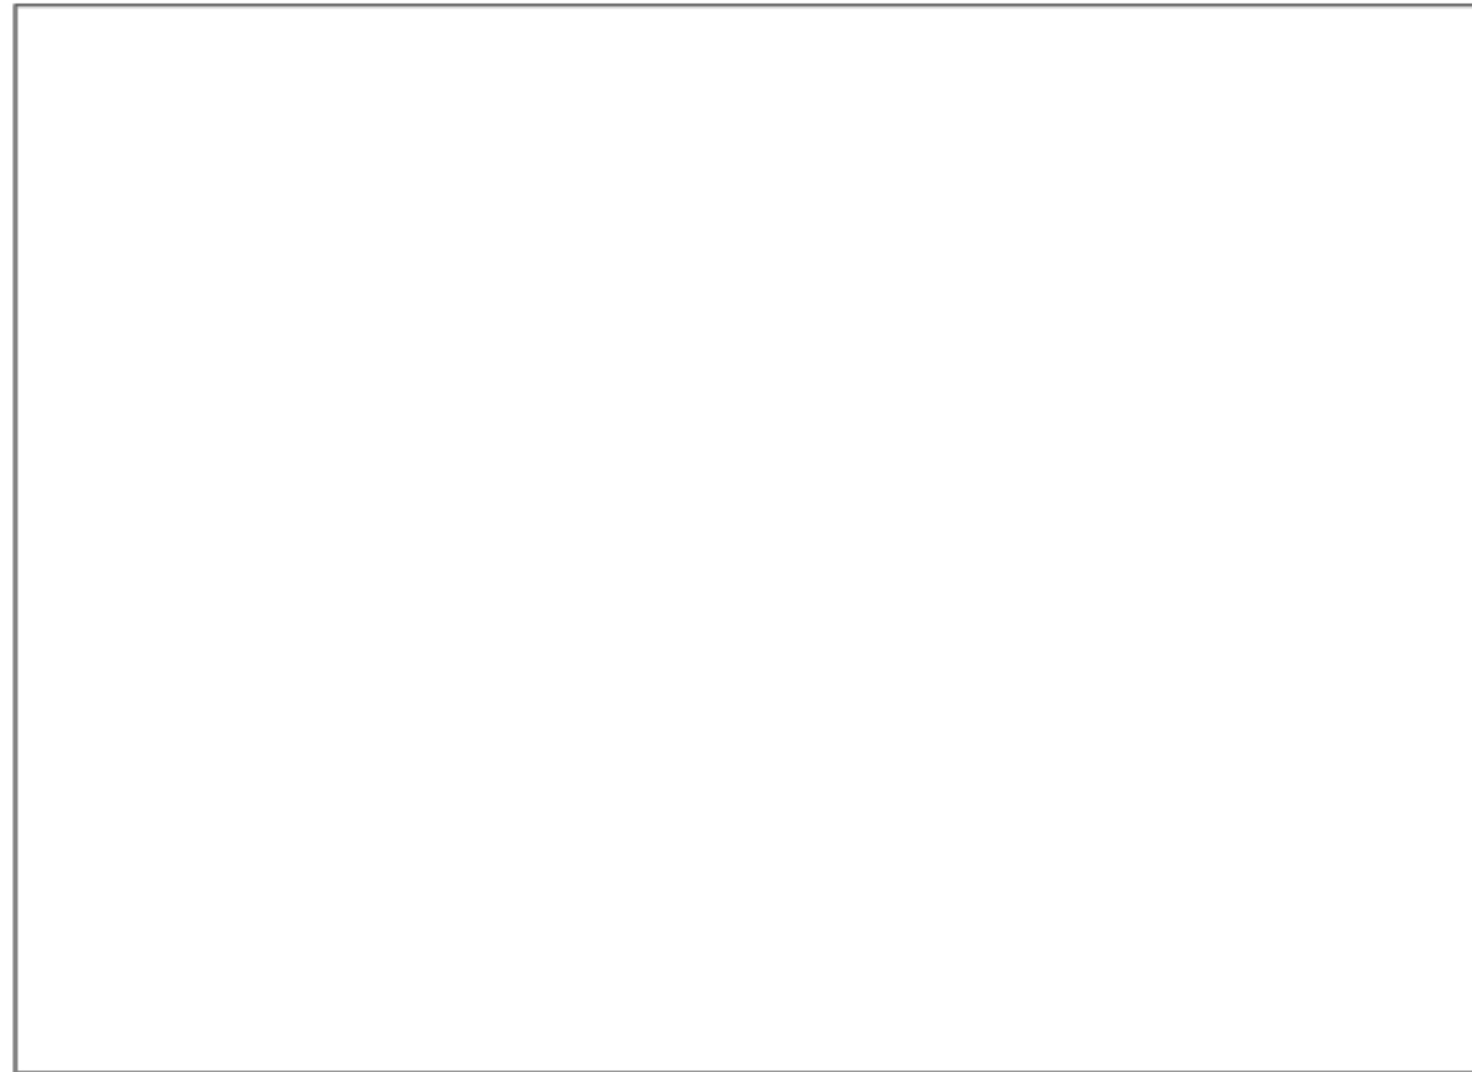
**Mark Peterson**

Senior Data Scientist, Alliance Data

# Probability thresholds

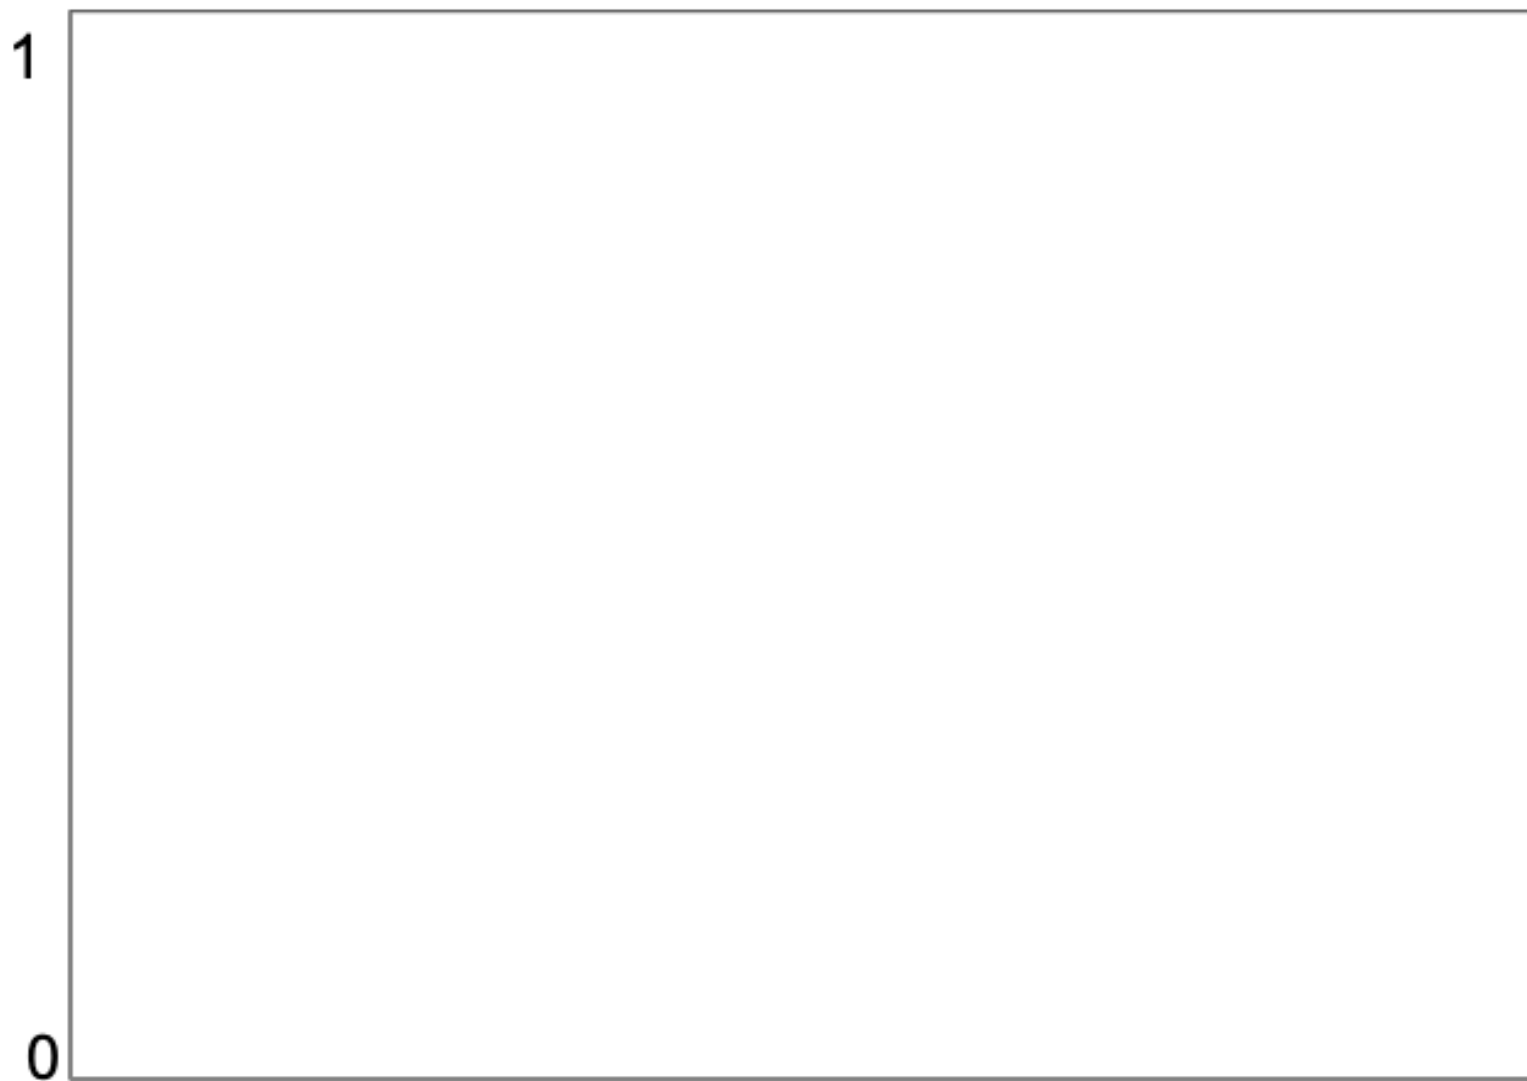
- Every prediction your classifier makes has an associated probability
  - Default probability threshold in scikit-learn: 50%
- 
- What if we vary this threshold?

## ROC Curve

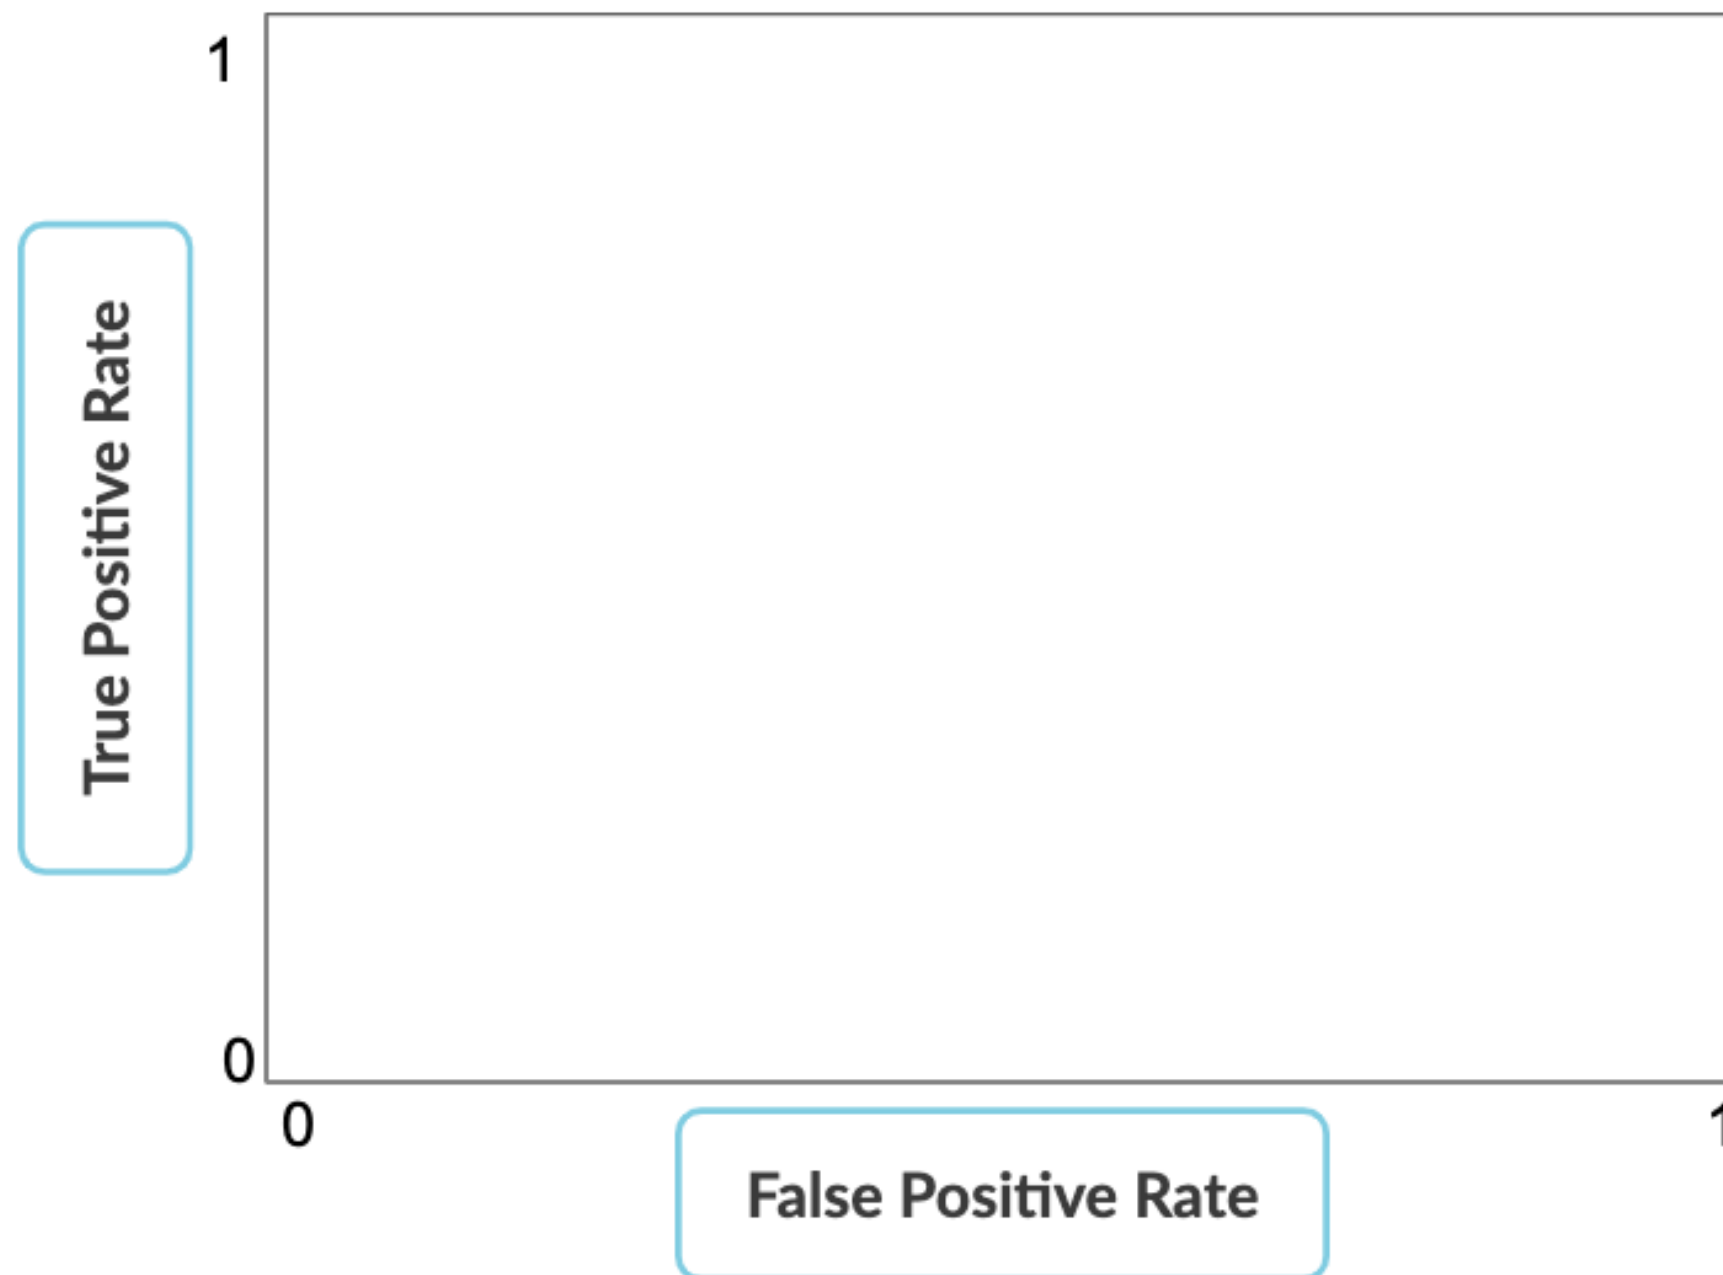


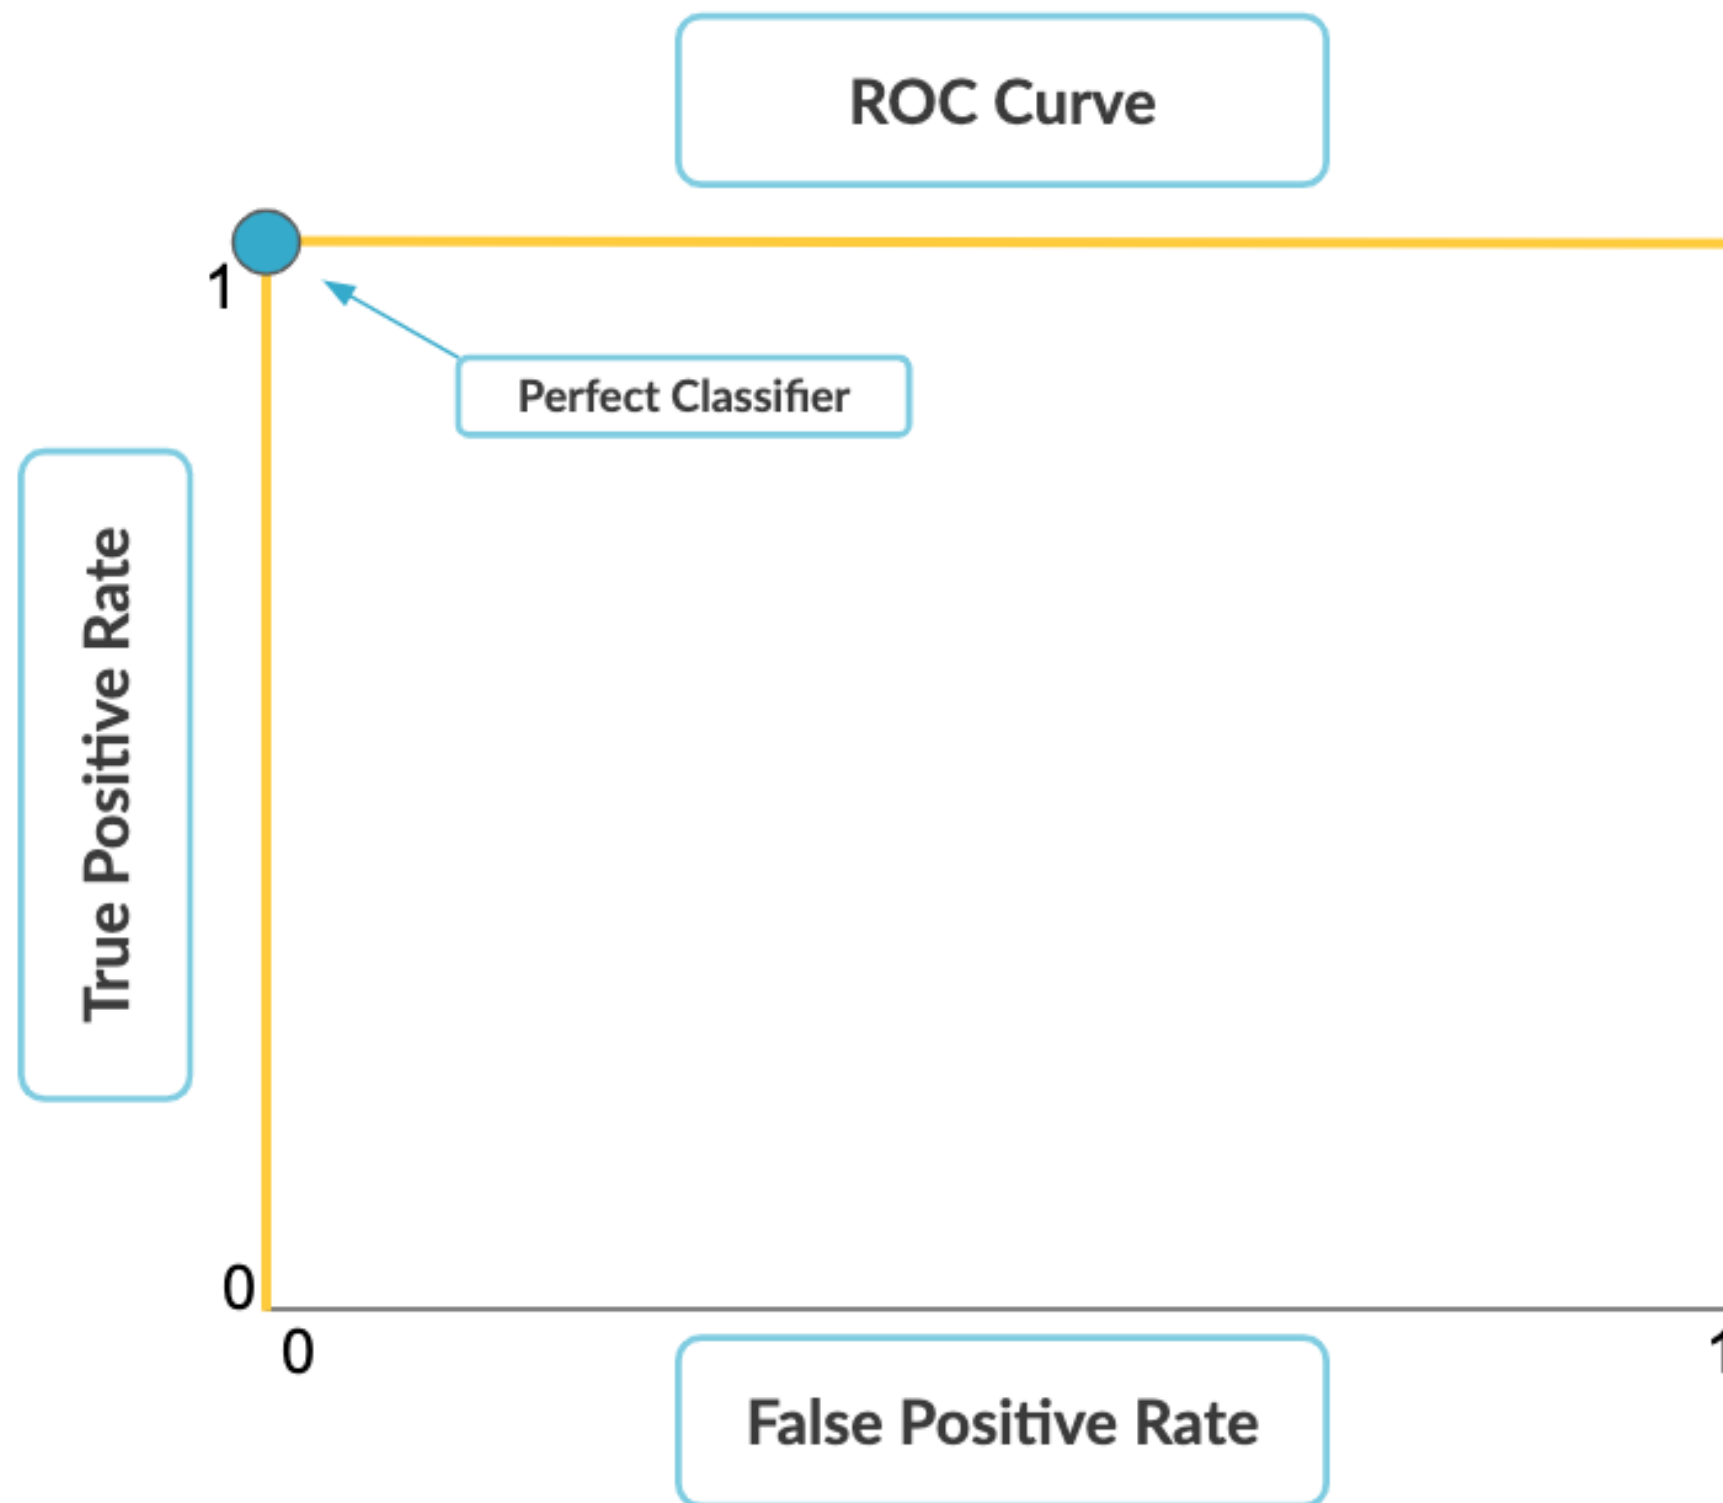
## ROC Curve

True Positive Rate



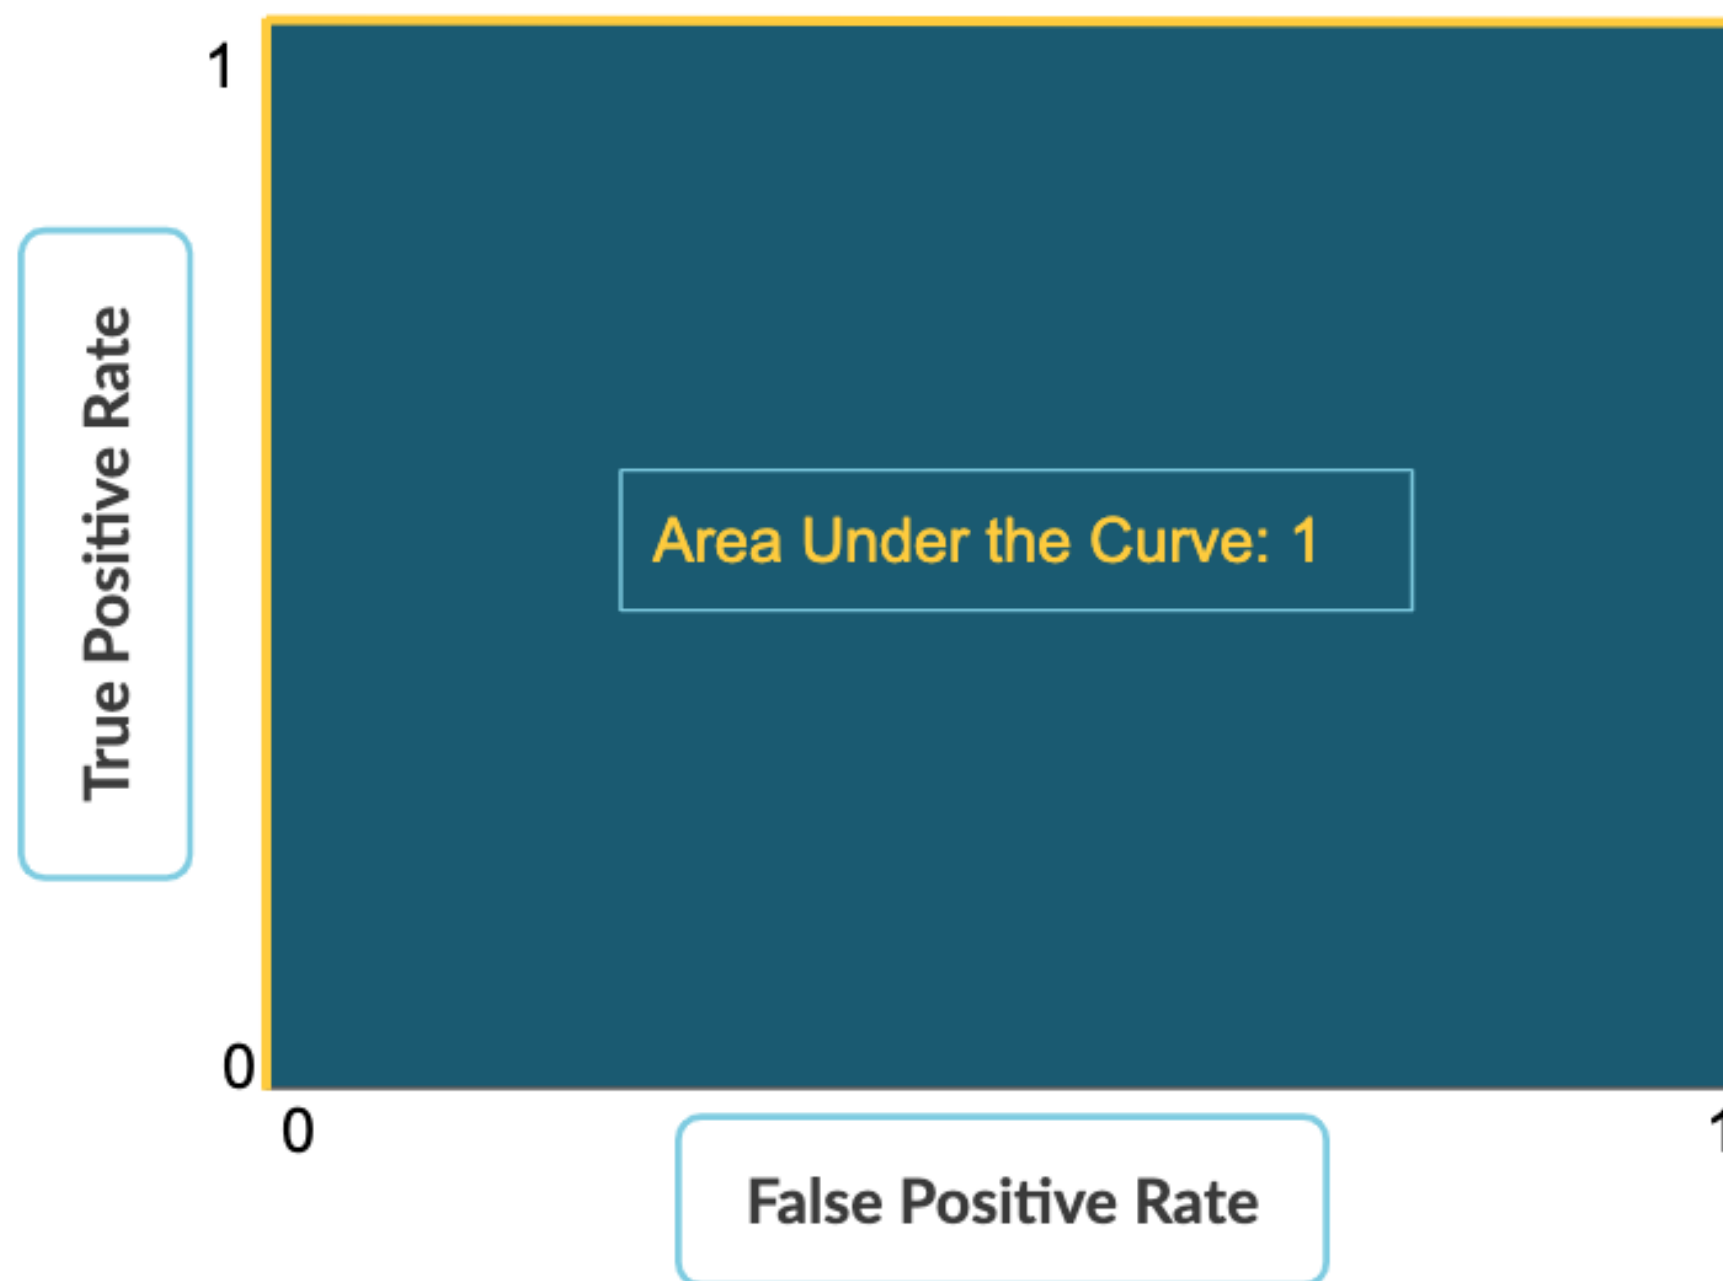
## ROC Curve



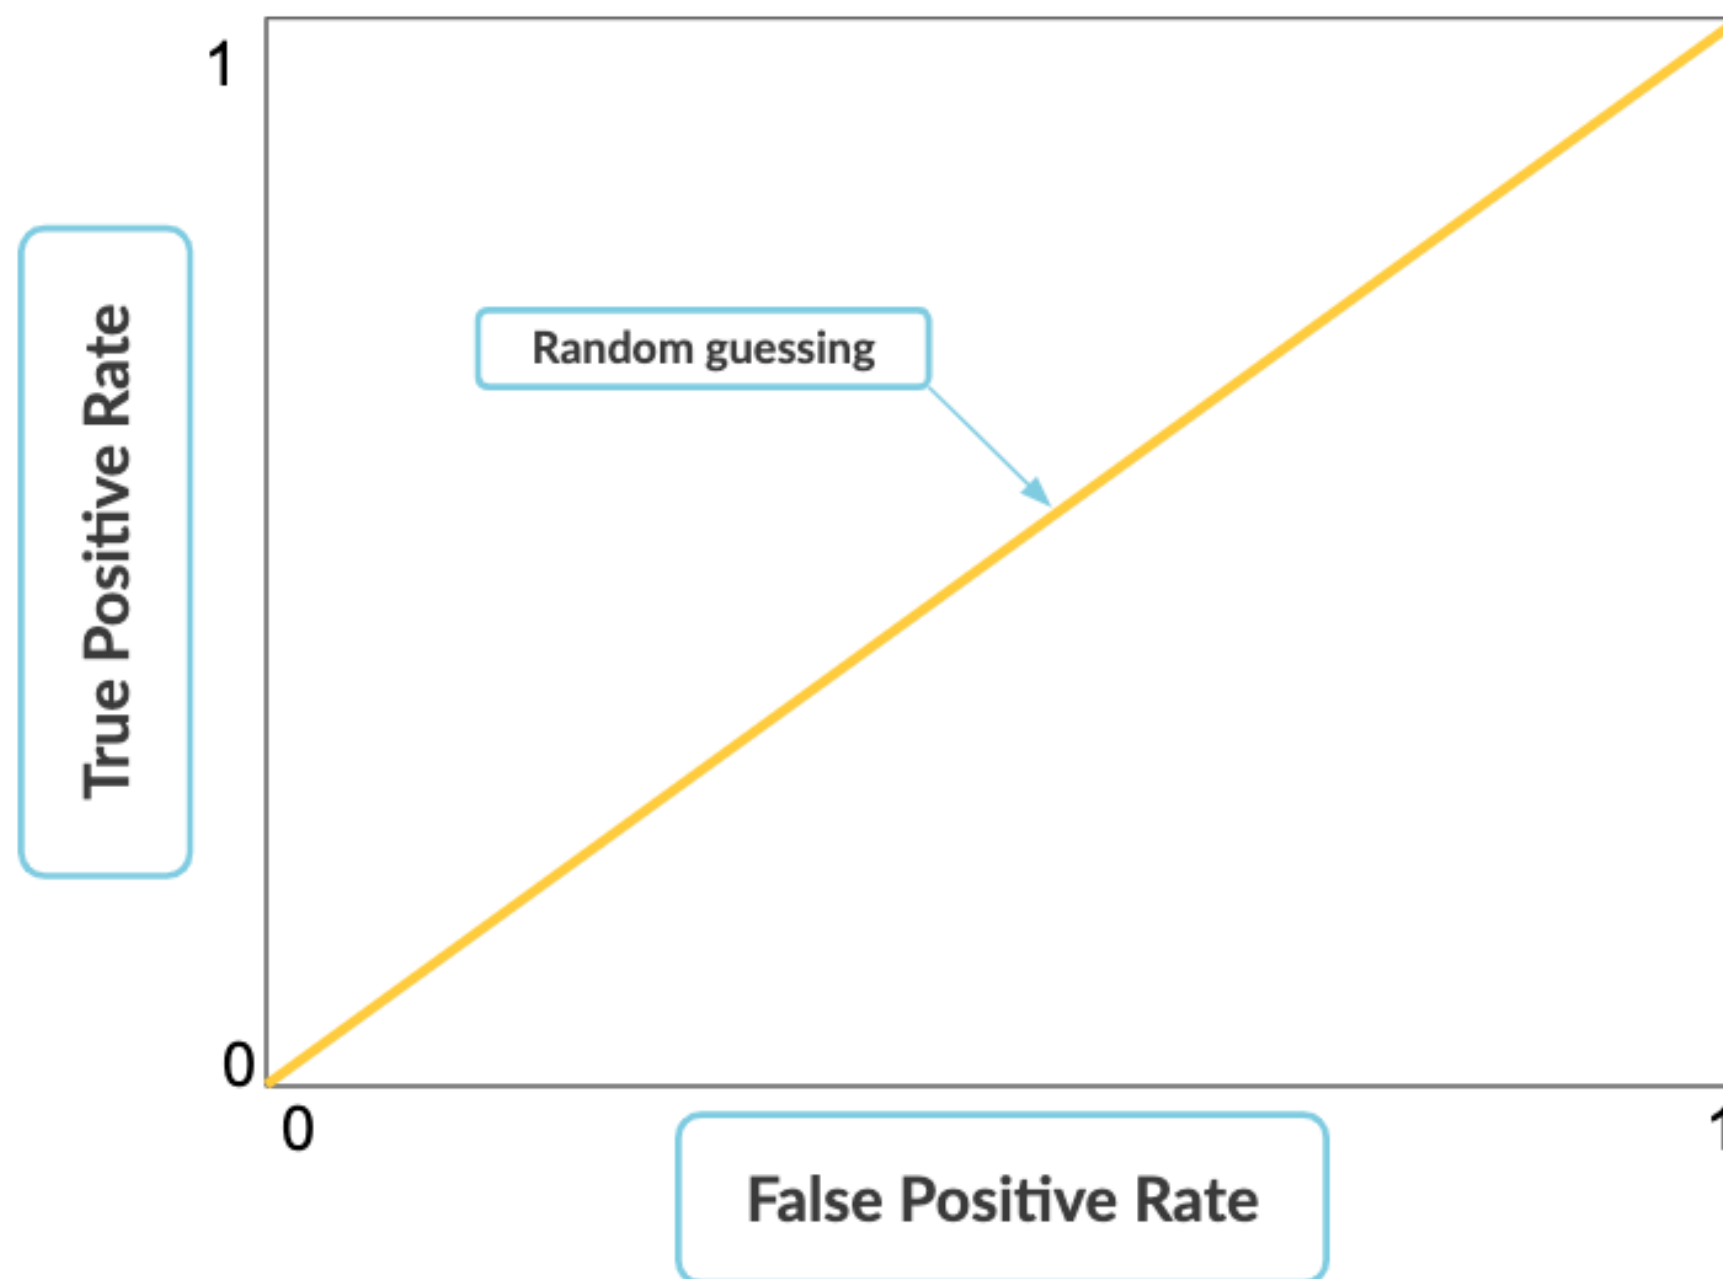




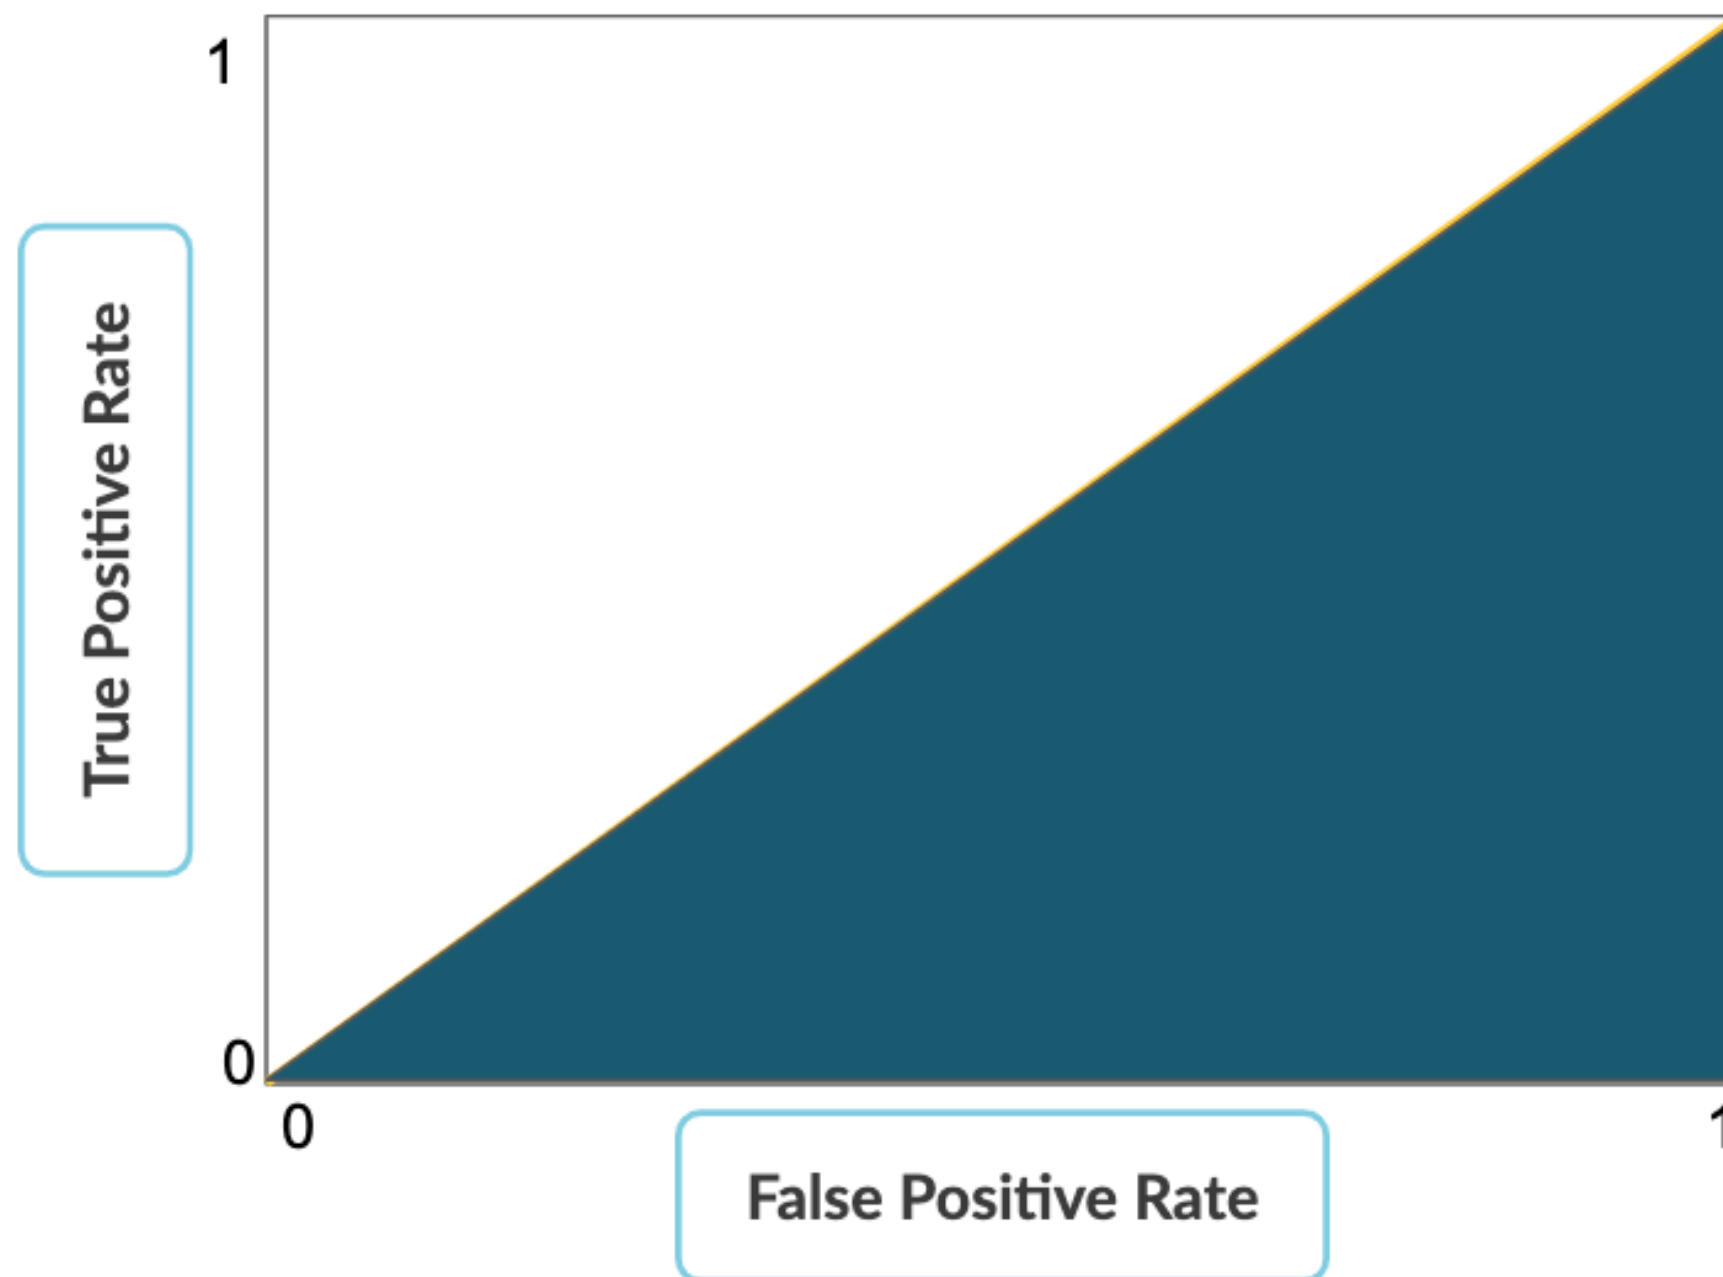
## ROC Curve



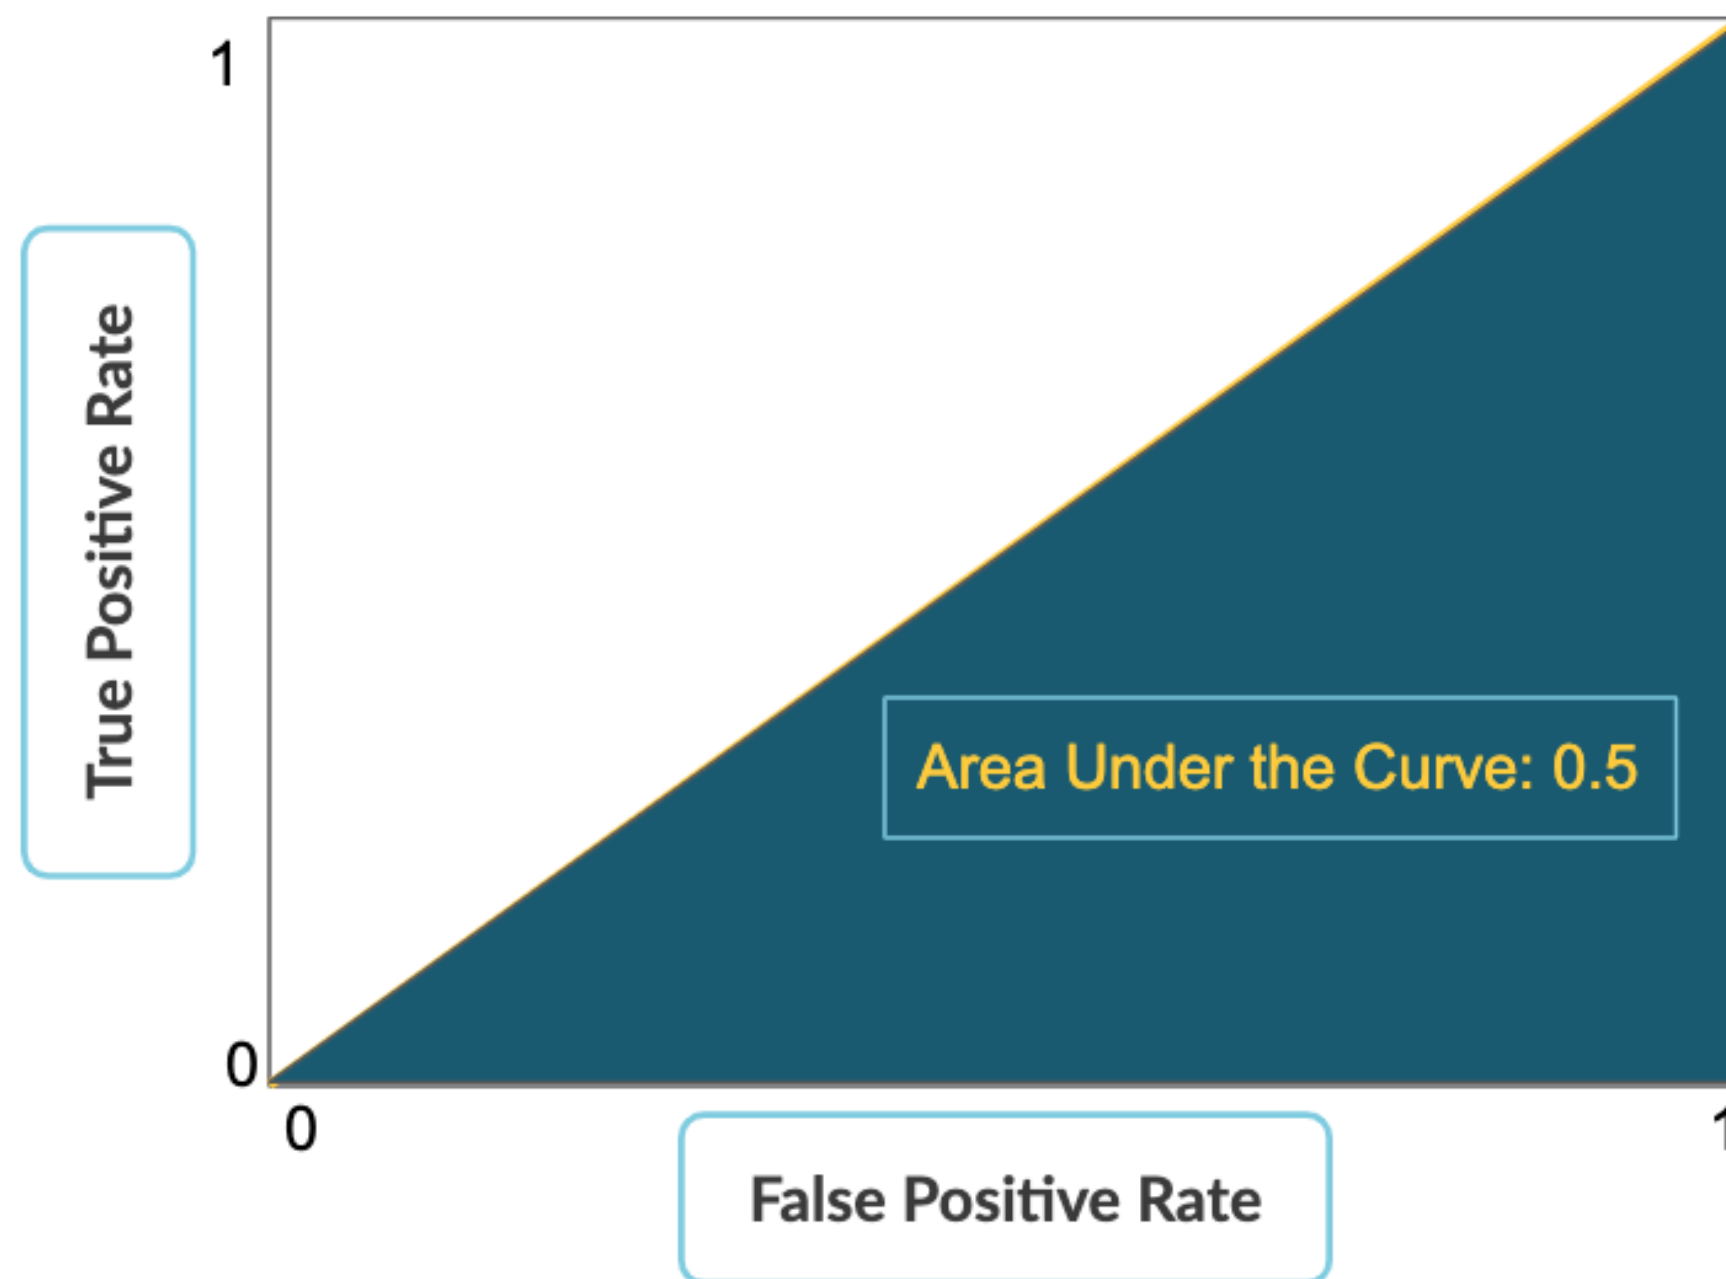
## ROC Curve



## ROC Curve



## ROC Curve



# Generating probabilities in sklearn

```
logreg.predict_proba(X_test)[: , 1]
```

```
array([[0.80188981, 0.19811019],  
       [0.96484075, 0.03515925],  
       [0.9182671 , 0.0817329 ],  
       ...,
```

```
y_pred_prob = logreg.predict_proba(X_test)[: , 1]
```

# ROC curve in sklearn

```
from sklearn.metrics import roc_curve

fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob)
```

```
import matplotlib.pyplot as plt

plt.plot(fpr, tpr)

plt.xlabel("False Positive Rate")

plt.ylabel("True Positive Rate")

plt.plot([0, 1], [0, 1], "k--")

plt.show()
```

# Area under the curve

```
from sklearn.metrics import roc_auc_score
```

```
auc = roc_auc_score(y_test, y_pred)
```

# Let's practice!

PREDICTING CUSTOMER CHURN IN PYTHON