Tuning your model

PREDICTING CUSTOMER CHURN IN PYTHON

Mark Peterson
Senior Data Scientist, Alliance Data

DataCamp

Refresher

```
from sklearn.svm import SVC

svc = SVC()

svc.fit(telco['data'], telco['target'])
```

```
SVC(C=1.0, cache_size=200, class_weight=None, coef0=0.0,
  decision_function_shape='ovr', degree=3, gamma='auto', kernel='rbf',
  max_iter=-1, probability=False, random_state=None, shrinking=True,
  tol=0.001, verbose=False)
```

Random forest hyperparameters

Parameter	Purpose
n_estimators	Number of trees
criterion	Quality of Split
max_features	Number of features for best split
max_depth	Max depth of tree
min_sample_splits	Minimum samples to split node
bootstrap	Whether Bootstrap samples are used

Grid search

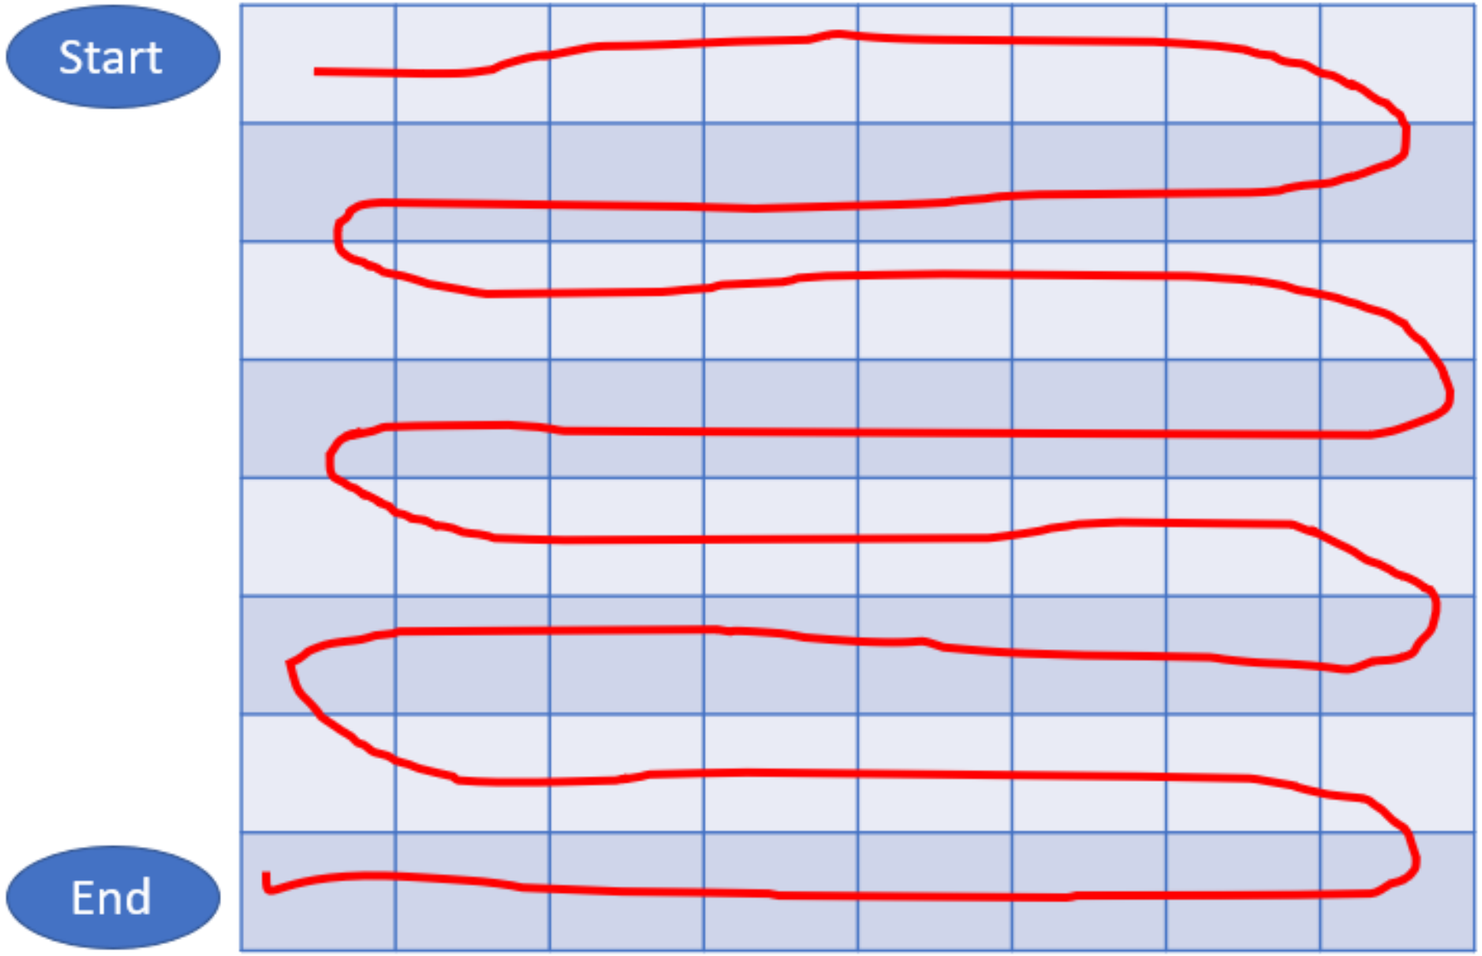

Grid search in sklearn

```
from sklearn.model_selection import GridSearchCV
param_grid = {'n_estimators': np.arange(10, 51)}
clf_cv = GridSearchCV(RandomForestClassifier(), param_grid)
clf_cv.fit(X, y)
clf_cv.best_params_
```

```
{'n_estimators': 43}
```

```
clf_cv.best_score_
```

```
0.9237923792379238
```

Happy tuning!

PREDICTING CUSTOMER CHURN IN PYTHON

DataCamp

Feature importances

PREDICTING CUSTOMER CHURN IN PYTHON

Mark Peterson
Senior Data Scientist, Alliance Data

DataCamp

Feature importances

- Scores representing how much each feature contributes to a prediction
- Effective way to communicate results to stakeholders

  - Which features are important drivers of churn?
  - Which features can be removed from the model?

Interpretability vs accuracy

- Different models have different strengths
- Need to balance prediction accuracy vs. interpretability

Random forest feature importances

```
random_forest = RandomForestClassifier()


random_forest.fit(X_train, y_train)
```

```
random_forest.feature_importances_
```

```
array([0.05069206, 0.04501006, 0.14427055, 0.08018487, 0.05222886,
       0.04418832, 0.11497537, 0.04463341, 0.12179754, 0.04756014,
       0.06818244, 0.05074536, 0.04616382, 0.03110577, 0.05826142])
```

Let's practice!

PREDICTING CUSTOMER CHURN IN PYTHON

DataCamp

Adding New Features

PREDICTING CUSTOMER CHURN IN PYTHON

Mark Peterson
Senior Data Scientist, Alliance Data

DataCamp

Additional Data Sources

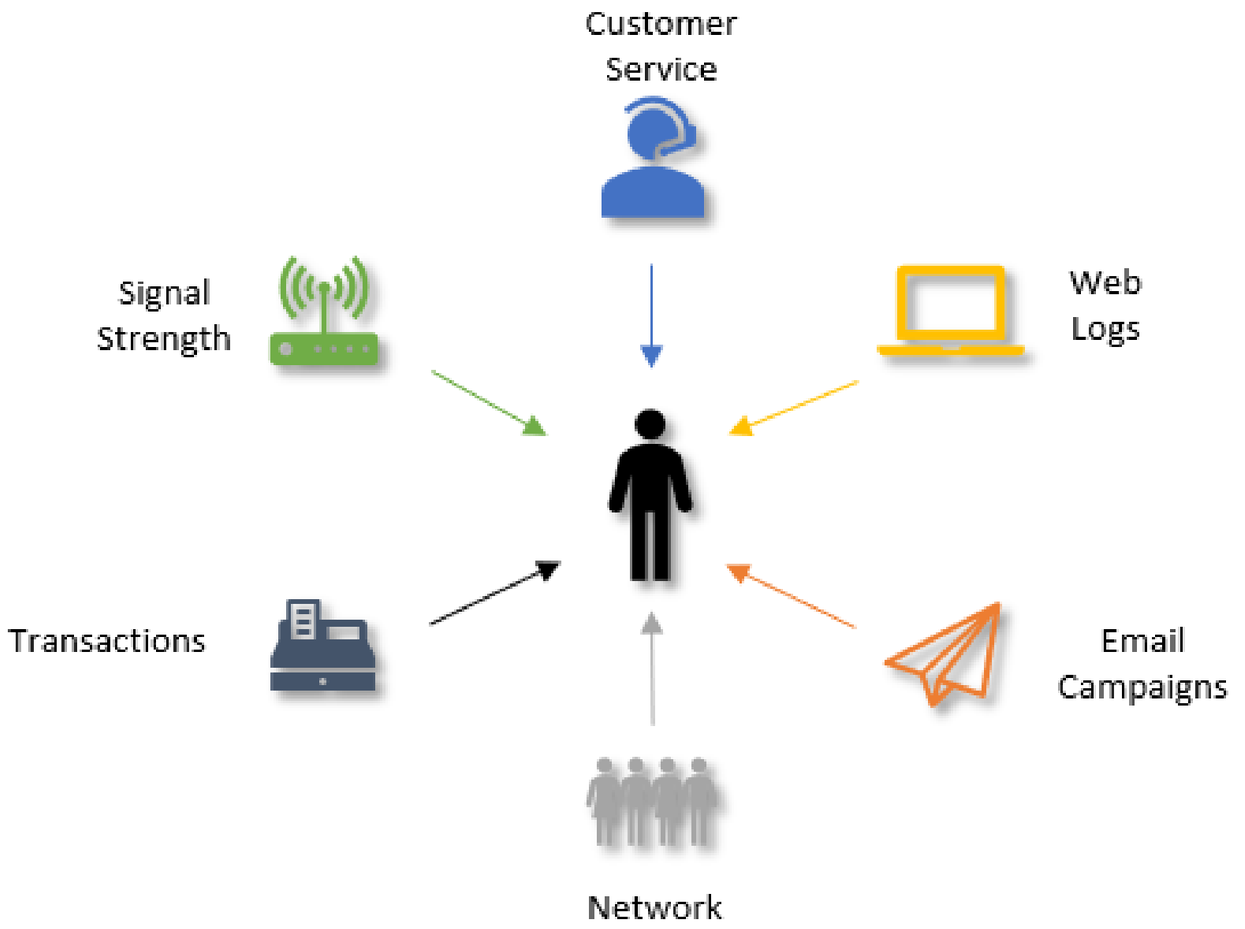

Churn Features

- Region Code
- Total Charges
- Total Minutes
- Minutes Per Call
- Cost Per Call
- Total Calls

Model Improvement

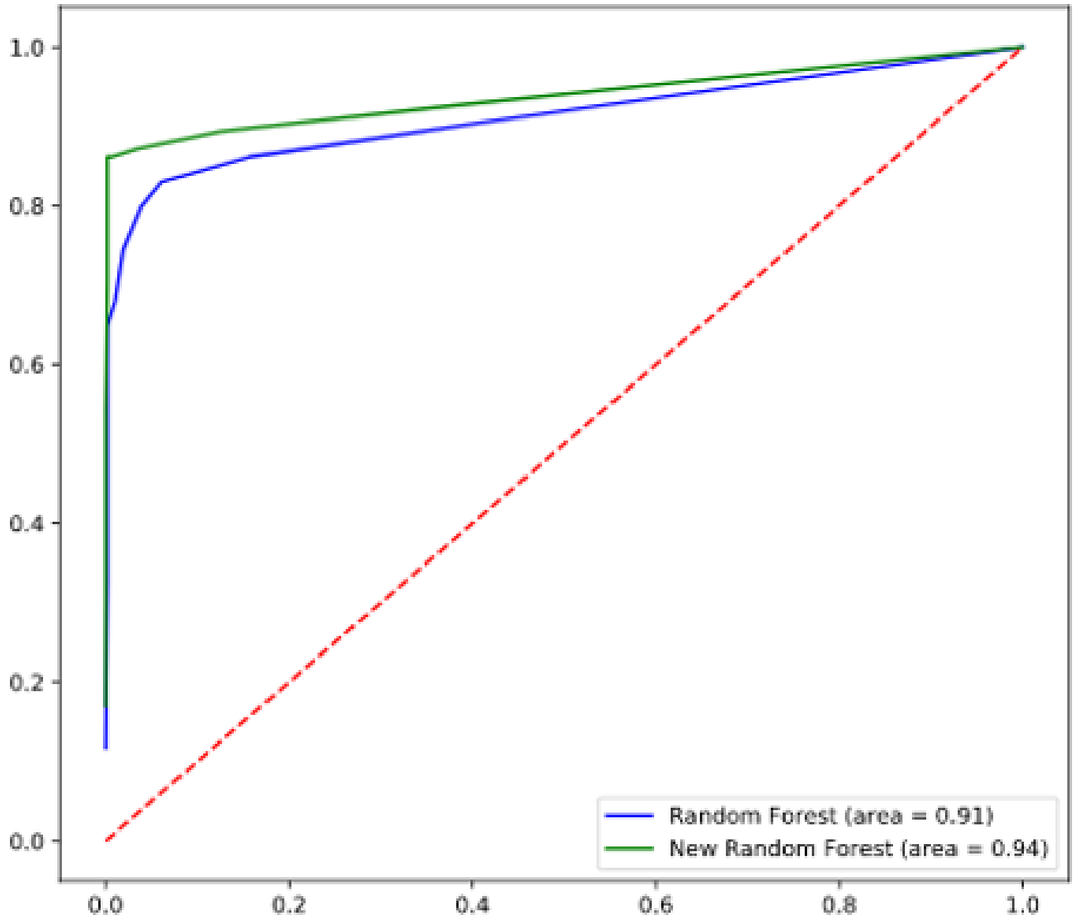

Benefits | Costs

- Benefits
  - Improved Return on Investment
  - Decreased Costs
  - Increased Performance
- Costs
  - Increased Complexity
  - Increased Resources
  - Increased Time to Operationalizing

Let's practice!

PREDICTING CUSTOMER CHURN IN PYTHON

DataCamp

Final thoughts

PREDICTING CUSTOMER CHURN IN PYTHON

Mark Peterson
Senior Data Scientist, Alliance Data

DataCamp

- Defined customer churn
- Exploratory data analysis

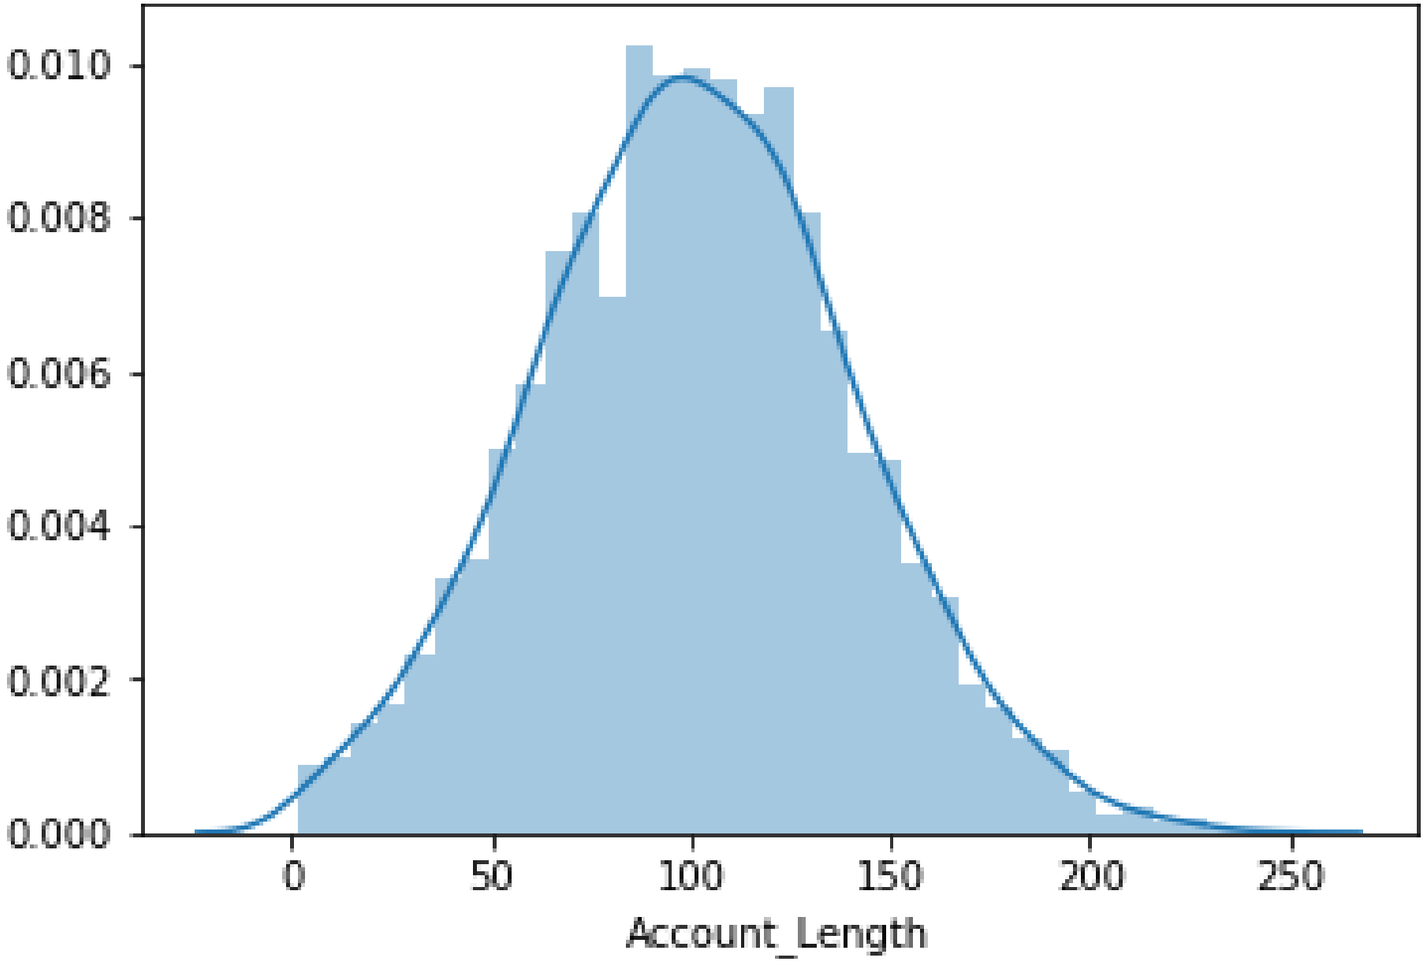

- Drop unnecessary features

State
KS
OH
NJ
OH

State_KS	State_OH	State_NJ
1	0	0
0	1	0
0	0	1
0	1	0

  - Feature scaling

- Making predictions
- Training and testing sets

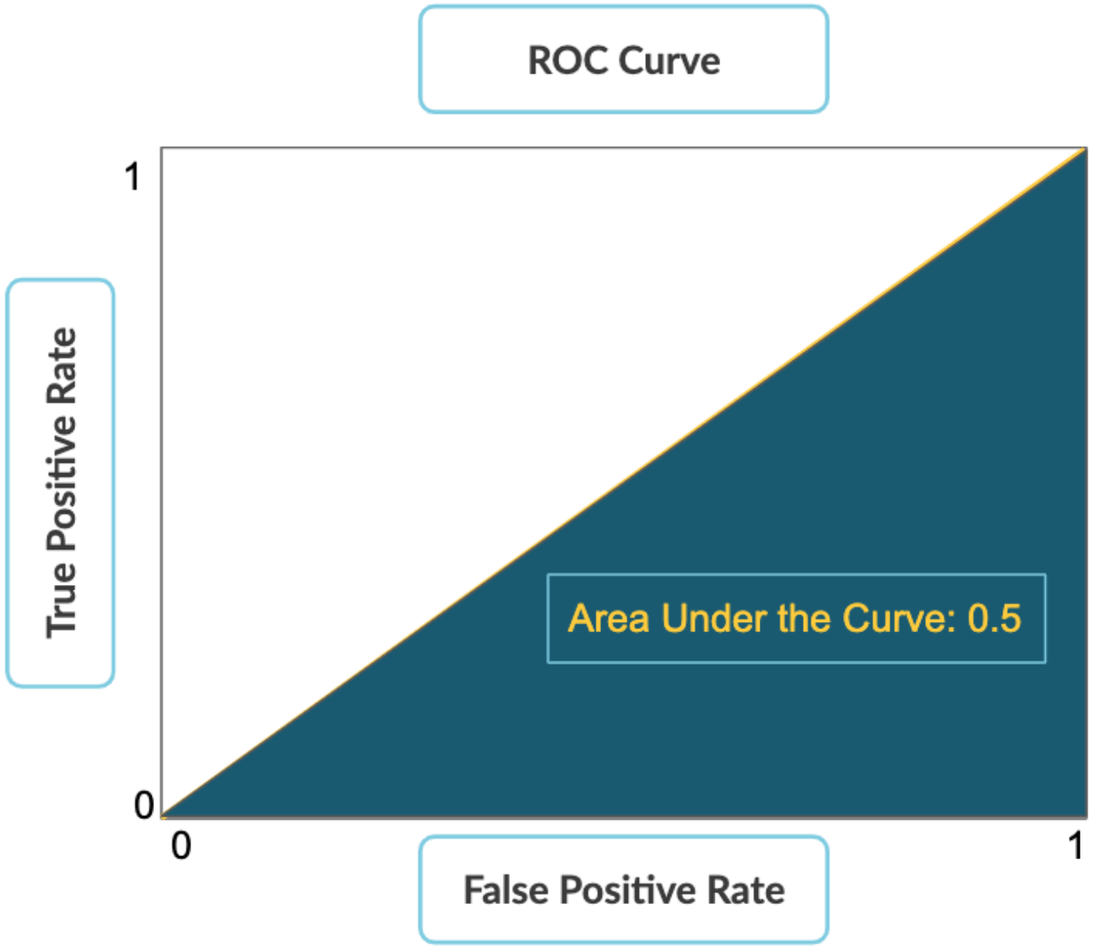

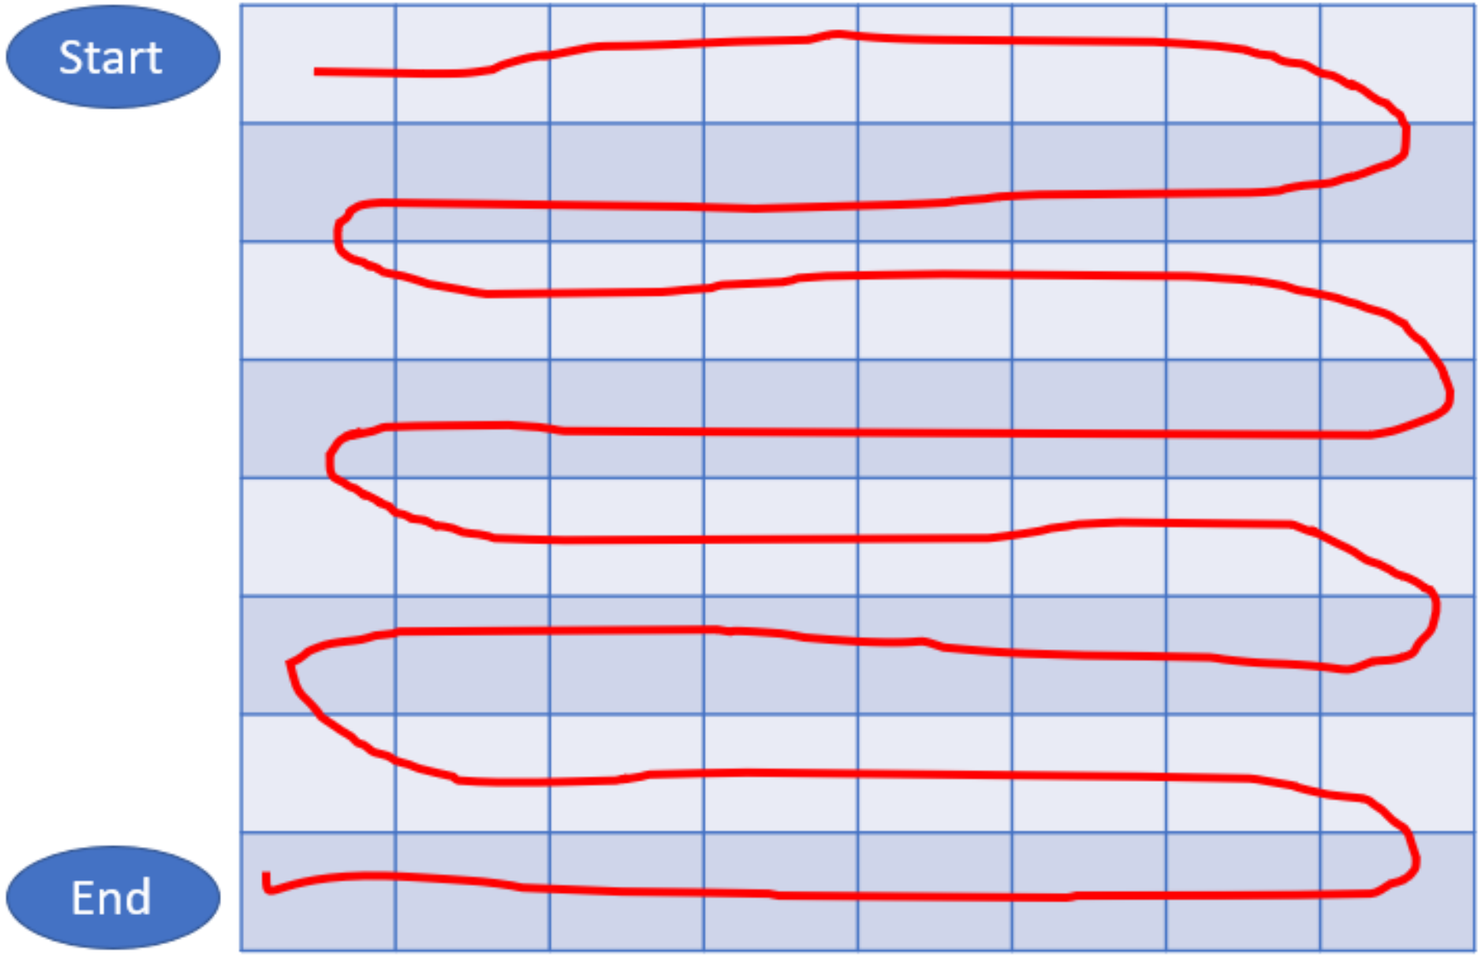
Start
End

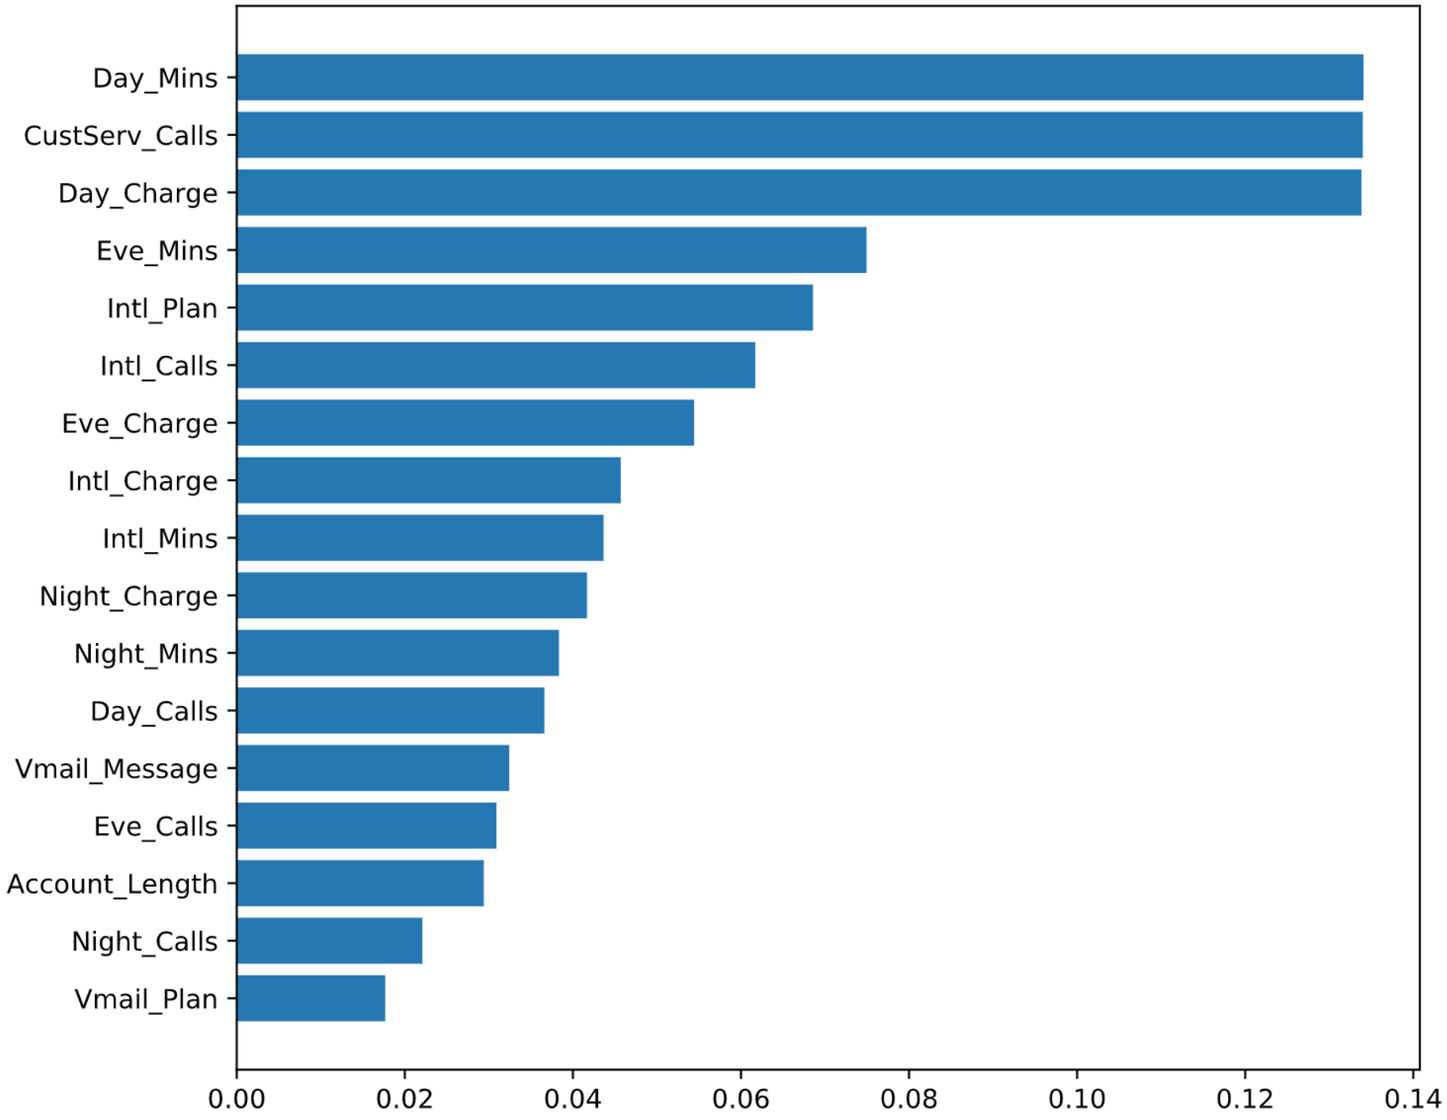
Day_Mins
CustServ_Calls
Day_Charge
Eve_Mins
Intl_Plan
Intl_Calls
Eve_Charge
Intl_Charge
Intl_Mins
Night_Charge
Night_Mins
Day_Calls
Vmail_Message
Eve_Calls
Account_Length
Night_Calls
Vmail_Plan
0.00
0.02
0.04
0.06
0.08
0.10
0.12
0.14

Churn Workflow

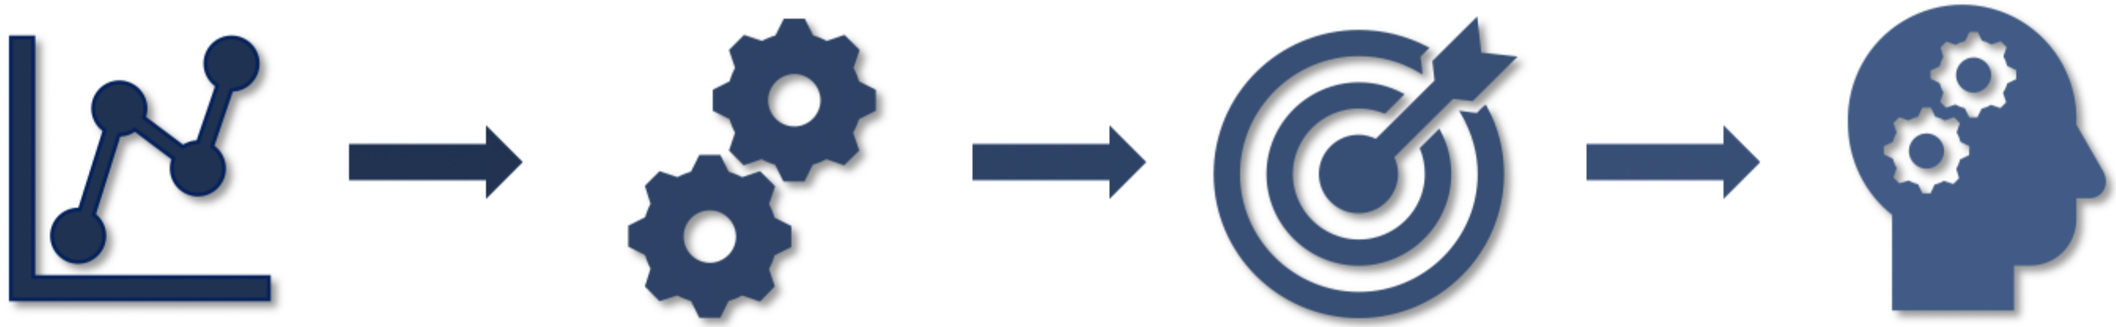

Recommended next steps

- Exploratory Data Analysis in Python
- Designing Machine Learning Workflows in Python

Additional resources

- Kaggle competitions
- Coursera advanced business analytics specialization

Great Work!

PREDICTING CUSTOMER CHURN IN PYTHON

DataCamp